

EDAV Final Project Report

Qi Feng, Zhida Zhang, Zhiyuan Zhao

INTRODUCTION

With the fierce competition in the labor pool nowadays, more and more high schoolers choose to pursue a college education after graduation. College education not only equips students with knowledge and skills that are essential for their development, but also expands their networks which could advance them in future career. Yet, the search for an appropriate college can be overwhelming, given the vast range of factors one has to consider and the limitation of knowledge one has for colleges at hand.

Therefore, in our project, we would like to provide a way for students and their families to compare across different colleges based on the **cost-value** tradeoffs catering to their own needs, academic and career goals. We split cost-value tradeoffs into multiple parts, and define **cost** as integration of *admission requirement* and *affordability* of the college, and **value** as the combination of *education quality*, *diversity*, *completion* and *earnings after graduation*. Through our project, we will conduct exploratory data analysis on each part mentioned.

To guide the project, we propose the main research question: **how the values are related to the costs?** Specifically, do schools with higher admission requirement and higher tuition fees provide higher education quality, diversity, completion and earnings?

Our team members include Qi Feng, Zhida Zhang, and Zhiyuan Zhao. Qi Feng and Zhiyuan Zhao are responsible for data quality, and Zhida Zhang is in charge of presenting the results on the interactive platform of Shiny. Other than those, we split the remaining tasks evenly by each taking on two parts of the cost-value tradeoffs to perform analysis as well as writing the report. All of our work can be found here in the Github repository: https://github.com/ZhangZhida/EDAV_FinalProject.

DESCRIPTION OF DATA

The dataset employed in this project was obtained from *The College Scorecard* (<https://collegescorecard.ed.gov/data/>), which integrates data from sources including IPEDS, NSLDS, and Department of Treasury. The data was collected from federal reporting of undergraduate degree-granting institutions of higher education. There were 1899 features in the raw dataset, covering topics including school, academics, admissions, student, cost, aid, repayment, completion and earnings. For our interest, we only focus on the data from 2014 to 2015, which is the most recent time span when earnings for graduates are available.

ANALYSIS OF DATA QUALITY

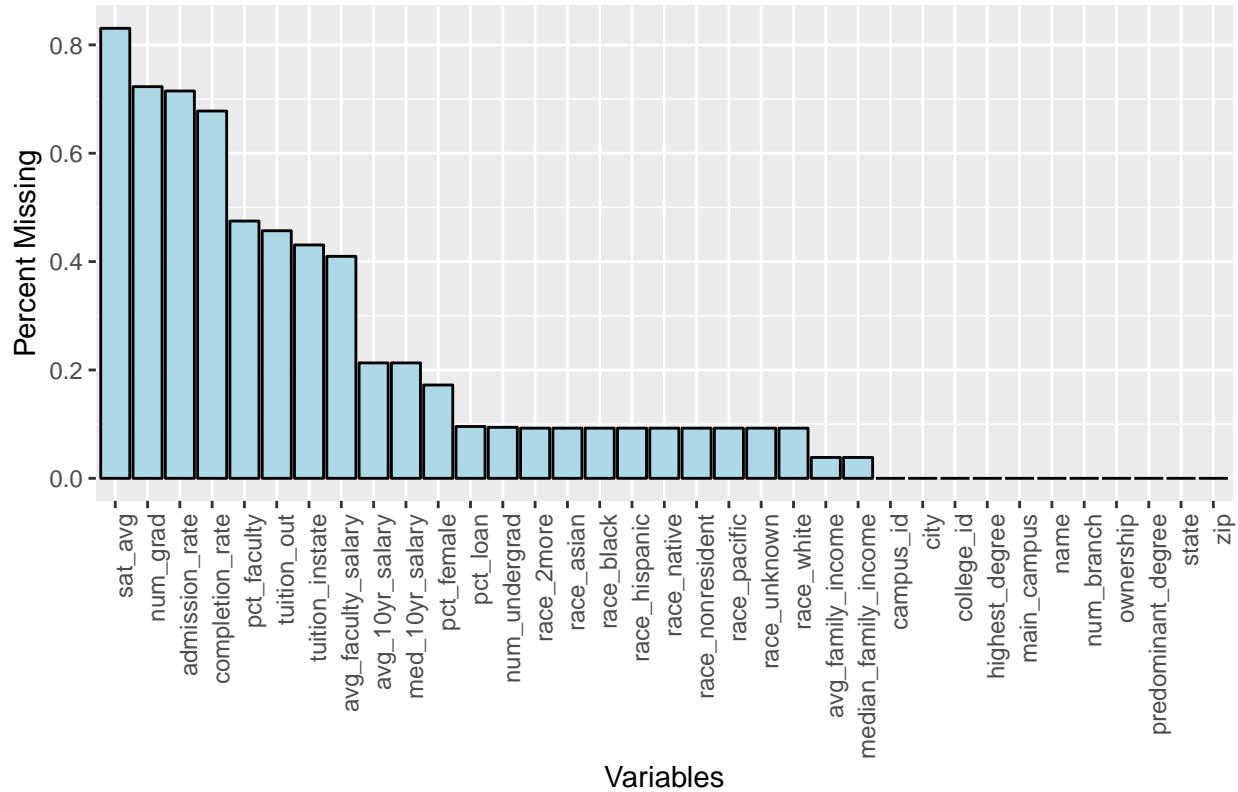
Since *The College Scorecard* project is supervised under the U.S. Department of Education, we are safe to presume the source of our data is valid. For our interest, we selected 35 features, of which 25 are continuous and 10 are categorical. A detailed data dictionary could be found at the GitHub repository: https://github.com/ZhangZhida/EDAV_FinalProject/blob/master/Data%20Dictionary.pdf. Data cleaning code can also be found at https://github.com/ZhangZhida/EDAV_FinalProject/blob/master/Data%20cleaning.R.

For missing value and patterns, we have the following analysis. According to the histogram below, for most schools, average SAT scores are missing. Number of graduate students, admission rate, and completion rate have missing values over 50%. Percent of full time faculty, in-state tuition, out-of-state tuition, and average faculty salary have missing values a little under 50%. Other variables have lower than 12% missing values.

```
library(tidyverse)
library(gridExtra)
college = read_rds("college.rds")
college[college=='NULL'] <- NA
a=data.frame(apply(college, 2, function(col)sum(is.na(col))/length(col)))
a$variables = rownames(a)
names(a) <- c("missing", "variables")
rownames(a) <- NULL

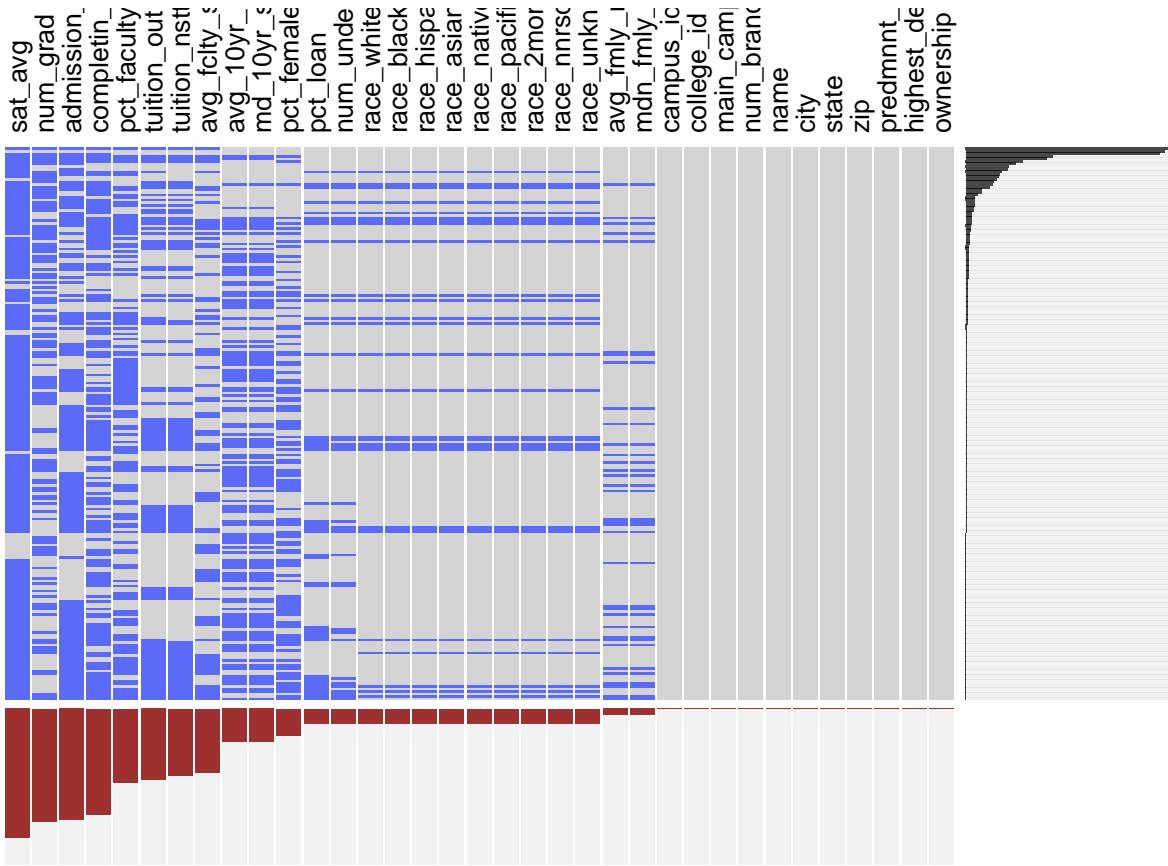
ggplot(a, aes(x=reorder(factor(a$variables), -a$missing), y=a$missing))+
  geom_bar(stat="identity", fill = "lightblue", color = "black")+
  ggtitle("Percent Missing by Variable")+
  labs(x = "Variables", y = "Percent Missing") +
  theme(axis.text.x = element_text(size = rel(1), angle = 90, hjust = 1),
        axis.text.y = element_text(size = rel(1)),
        strip.text = element_text(size = rel(1)))
```

Percent Missing by Variable



The most frequent missing pattern is missing average SAT, number of graduate, admission rate, completion rate, percent of full time faculty, in-state tuition, out-of-state tuition, and average faculty salary at the same time. The eight variables are also the ones with the most missing values. For race, the percentages for missing values are the same among different races, and the missing pattern for race percent is the same across the data, meaning when one race percent is missing, all are missing.

```
library(extractat)
visna(college, sort="b")
```



MAIN ANALYSIS

COST

Admission Requirement

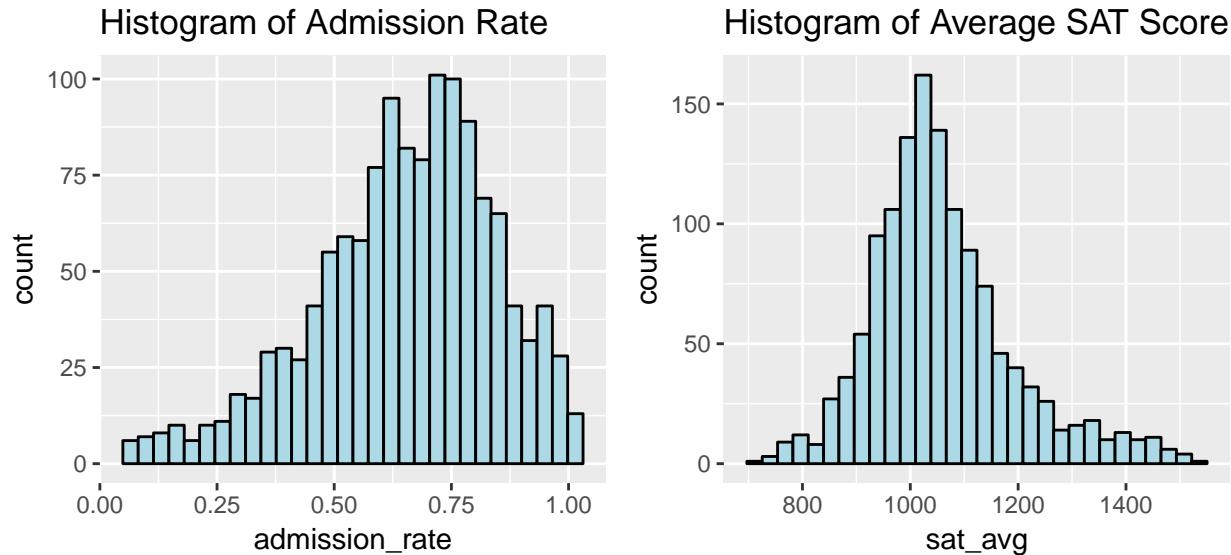
To get a sense of how admission and SAT scores are distributed respectively, we plot the histogram on these two variables. From the histogram of admission rate, we could see that most schools have the admission rate range between 0.5 to 0.85. Only a small proportion of schools have admission rate below 0.25. From the SAT average histogram, we see that the majority of students enrolled into universities have scores in the range from 900 to 1200. The higher the score is, the fewer students are admitted.

```
library(ggplot2)
college <- readRDS('college.rds')
college_no_na <- college[with(college, (!is.na(sat_avg)) & (!is.na(admission_rate))), ]

plot_admi <- ggplot(college_no_na, aes(x = admission_rate)) +
  geom_histogram(fill = "lightblue", color = "black") +
  ggtitle("Histogram of Admission Rate")

plot_sat <- ggplot(college_no_na, aes(x = sat_avg)) +
  geom_histogram(fill = "lightblue", color = "black") +
  ggtitle("Histogram of Average SAT Score")

grid.arrange(plot_admi, plot_sat, ncol = 2)
```



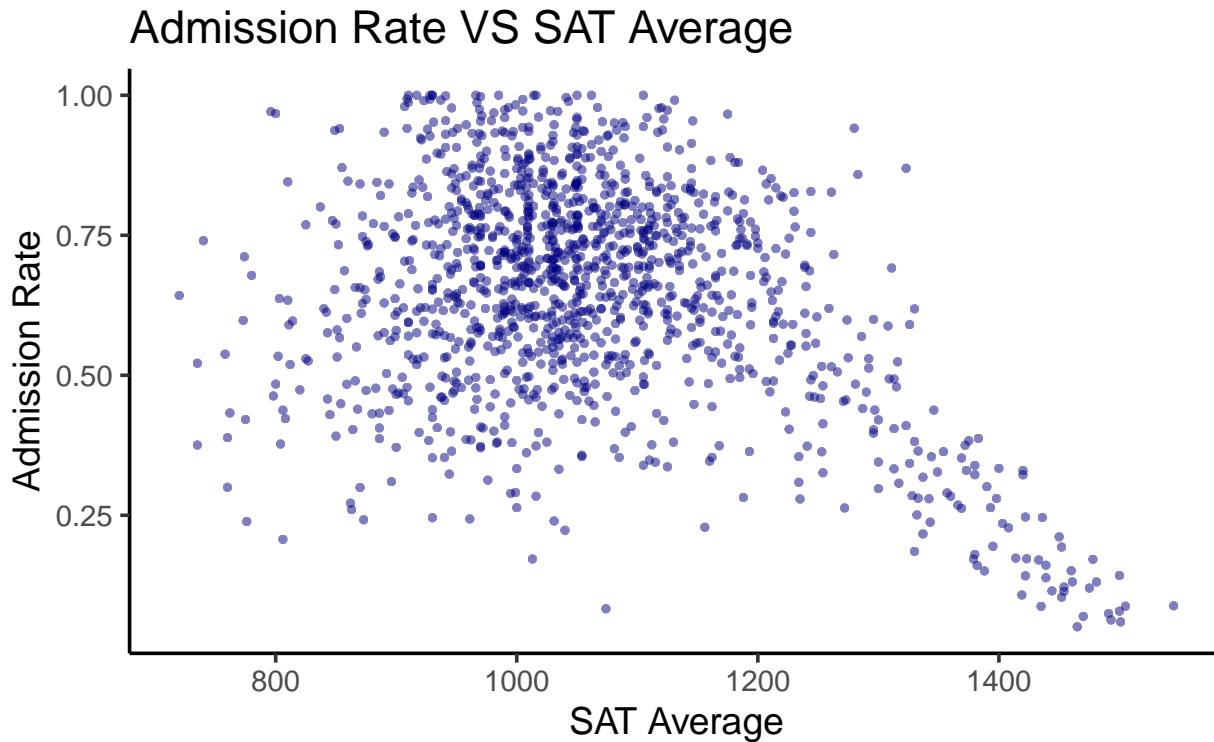
However, is there a relationship between admission rate and average SAT score? Intuitively, we should expect that higher SAT score goes along with lower admission rate. Utilizing a scatterplot, we can see that there are mainly two patterns. For those schools with SAT average below 1200, admission rate seems not to be closely related to SAT average score. That is, SAT average score does not appear to affect admission. On the other hand, for those universities with SAT average higher than 1200, we could clearly see a decreasing trend of admission rate with increasing SAT average. Such a result confirms our intuition that when a school has a high standard for SAT score, it tends to admit less students.

```
college_no_na <- college[with(college, (!is.na(sat_avg)) & (!is.na(admission_rate))), ]
sat_avg <- college_no_na$sat_avg
admission_rate <- college_no_na$admission_rate
```

```

baseplt <- ggplot(college_no_na, aes(x = sat_avg, y = admission_rate))
baseplt + geom_point(alpha= 0.5, color = "navy", stroke = 0) +
  theme_classic(14) +
  labs(title = "Admission Rate VS SAT Average",
       x = "SAT Average",
       y = "Admission Rate")

```



We then plot the density contour lines based on the scatter plots as a follow-up for our observation. From the left-hand side graph below, we find a sharp declining pattern: as SAT average becomes slightly higher, the admission rate drops a lot. The reason for this may be that scores higher than 1400 are very rare. Meanwhile, we see some universities with relatively lower SAT average still have very low admission rates. This might be due to the fact that SAT score is only one single requirement for these universities, so that they may have more emphasis on parts other than SAT score, such as internship experience, sports talents, etc.

```

college_no_na_low <- college[with(college, (!is.na(sat_avg)) &
                                    (!is.na(admission_rate)) & (admission_rate < 0.25)), ]
college_no_na_high <- college[with(college, (!is.na(sat_avg)) &
                                    (!is.na(admission_rate)) & (admission_rate >= 0.25)), ]
sat_avg <- college_no_na$sat_avg
admission_rate <- college_no_na$admission_rate

plot_low <- ggplot(college_no_na_low, aes(x = sat_avg, y = admission_rate))+
  geom_point(alpha= 0.5, color = "navy", stroke = 0) +
  theme_classic(14) + theme(plot.title = element_text(size = rel(0.8))) +
  geom_density_2d(color = "red", bins = 5) +
  labs(title = "Low Admission Rate Plot",
       x = "SAT Average",
       y = "Admission Rate [0, 0.25]")

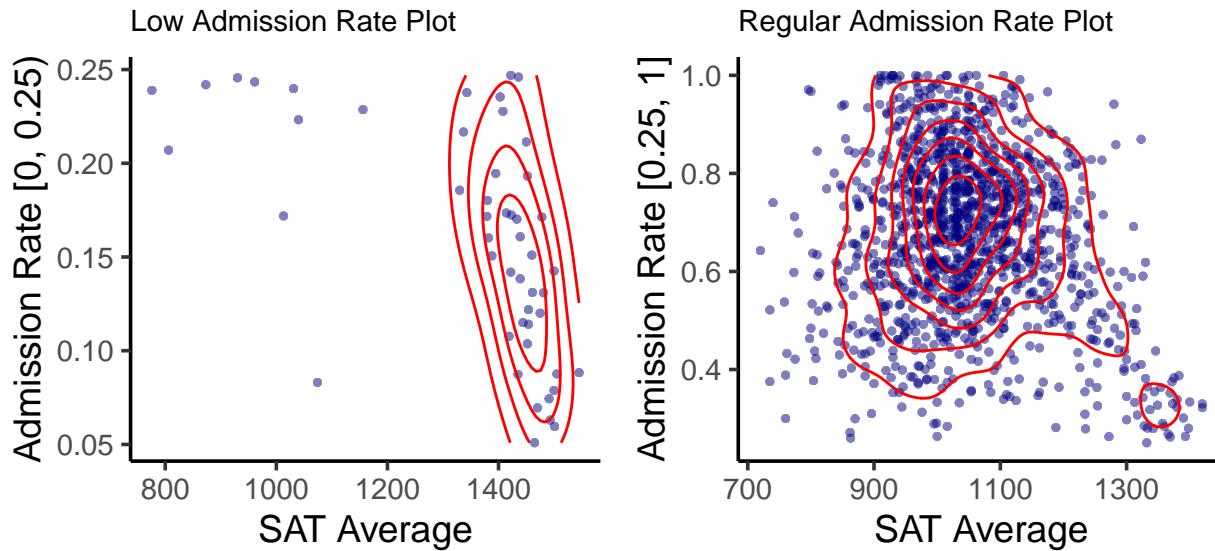
plot_high <- ggplot(college_no_na_high, aes(x = sat_avg, y = admission_rate))+

```

```

geom_point(alpha= 0.5, color = "navy", stroke = 0) +
theme_classic(14) + theme(plot.title = element_text(size = rel(0.8))) +
geom_density_2d(color = "red", bins = 8) +
labs(title = "Regular Admission Rate Plot",
x = "SAT Average",
y = "Admission Rate [0.25, 1]")
grid.arrange(plot_low, plot_high, ncol = 2)

```



Affordability

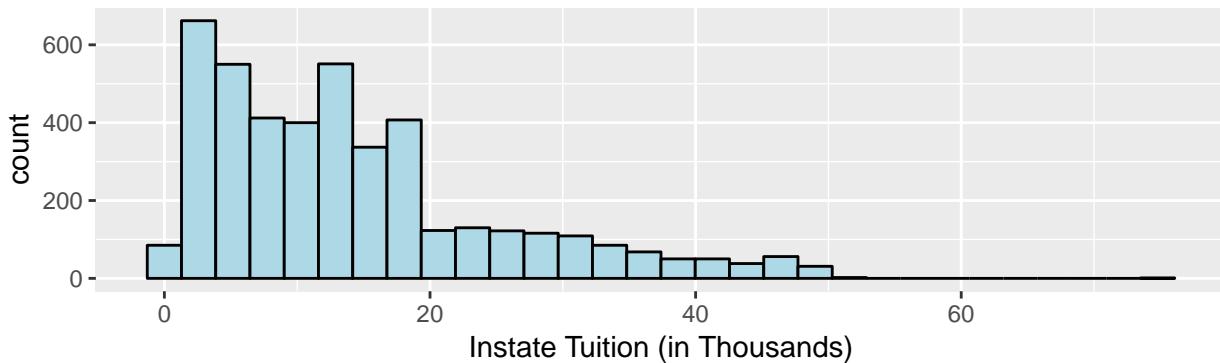
Instate and out-of-state tuition are the tuition and fees for instate and out-of-state students, respectively. From the histogram below, we can see that the distribution of the two tuitions are similar, and they both skewed to the right. Out-of-state students have overall higher tuition fee than instate students.

```

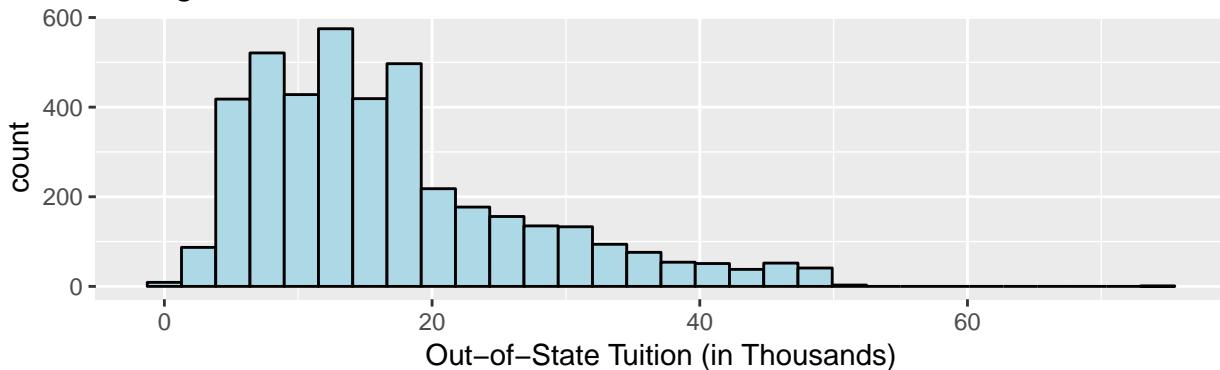
college = read_rds("college.rds")
college[college=="NULL"] <- NA
college1 = college[!is.na(college$tuition_instate), ]
college2 = college[!is.na(college$tuition_out), ]
college1$tuition_instate = college1$tuition_instate/1000
college2$tuition_out = college2$tuition_out/1000
recode <- c('1'='public','2'='nonprofit','3'='profit')
college1$ownership <- recode[college1$ownership]
college1$ownership<-fct_relevel(college1$ownership,'public','nonprofit','profit')
college2$ownership <- recode[college2$ownership]
college2$ownership<-fct_relevel(college2$ownership,'public','nonprofit','profit')
p3<-ggplot(college1, aes(x=tuition_instate))+
  geom_histogram(bins = 30,fill = "lightblue", color = "black")+
  labs(x="Instate Tuition (in Thousands)")+
  ggtitle("Histogram of Instate Tuition")
p4<-ggplot(college2, aes(x=tuition_out))+
  geom_histogram(bins = 30,fill = "lightblue", color = "black")+
  labs(x="Out-of-State Tuition (in Thousands)")+
  ggtitle("Histogram of Out-of-State Tuition")
grid.arrange(p3,p4,ncol=1)

```

Histogram of Instate Tuition



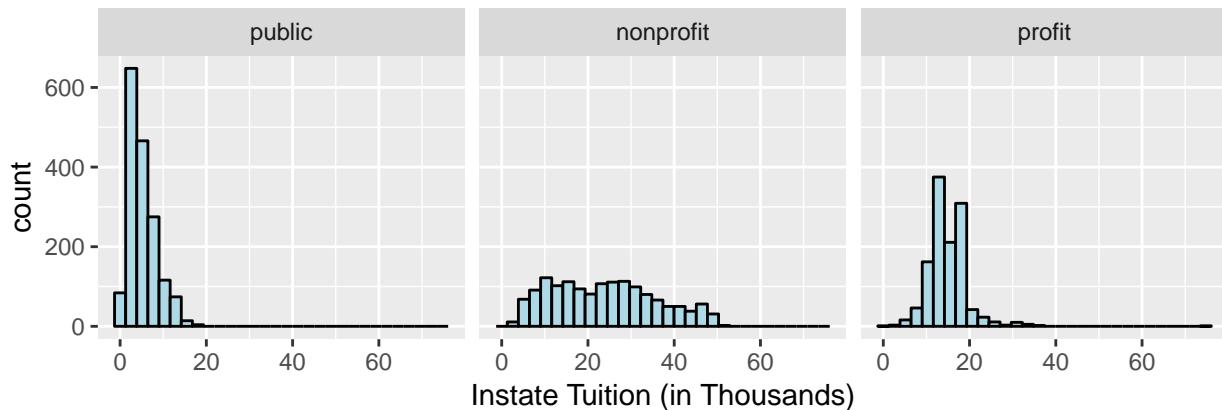
Histogram of Out-of-State Tuition



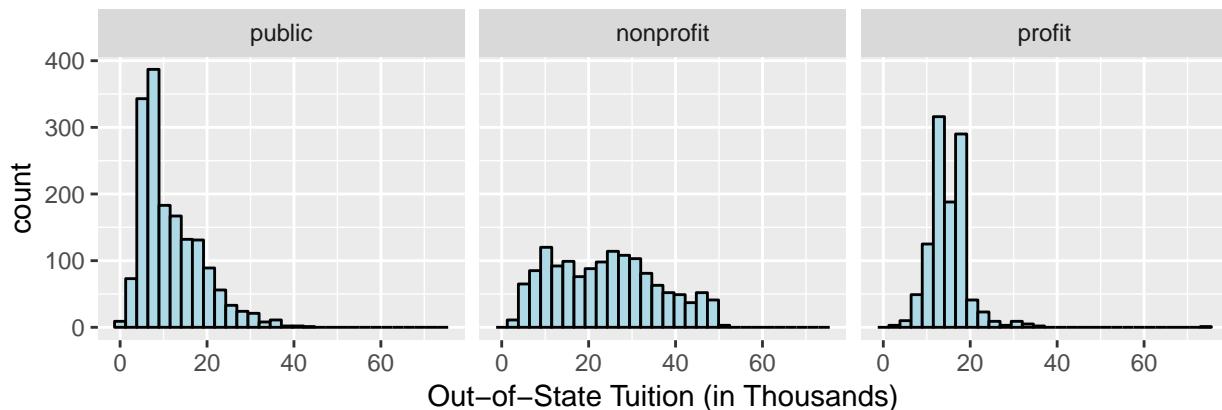
By faceting instate and out-of-state tuition into different ownership groups, we can see the different distributions for public, private nonprofit, and private for-profit schools, but the patterns for instate and out-of-state tuition are similar. For public schools, both instate and out-of-state tuition are the lowest among all three ownership types, and the distribution is unimodal and skewed to the right. For private nonprofit schools, tuition has a larger range with a bimodal distribution. Therefore, tuition differs for students from schools of different ownership types. We would recommend public schools for students with limited tuition budget.

```
p5 <- ggplot(college1, aes(x=tuition_instate))+
  geom_histogram(bins = 30, fill = "lightblue", color = "black")+
  facet_grid(.~ownership)+
  labs(x="Instate Tuition (in Thousands)")+
  ggtitle("Histogram of Instate Tuition by Ownership")
p6 <- ggplot(college2, aes(x=tuition_out))+
  geom_histogram(bins = 30, fill = "lightblue", color = "black")+
  facet_grid(.~ownership)+
  labs(x="Out-of-State Tuition (in Thousands)")+
  ggtitle("Histogram of Out-of-State Tuition by Ownership")
grid.arrange(p5,p6,ncol=1)
```

Histogram of Instate Tuition by Ownership

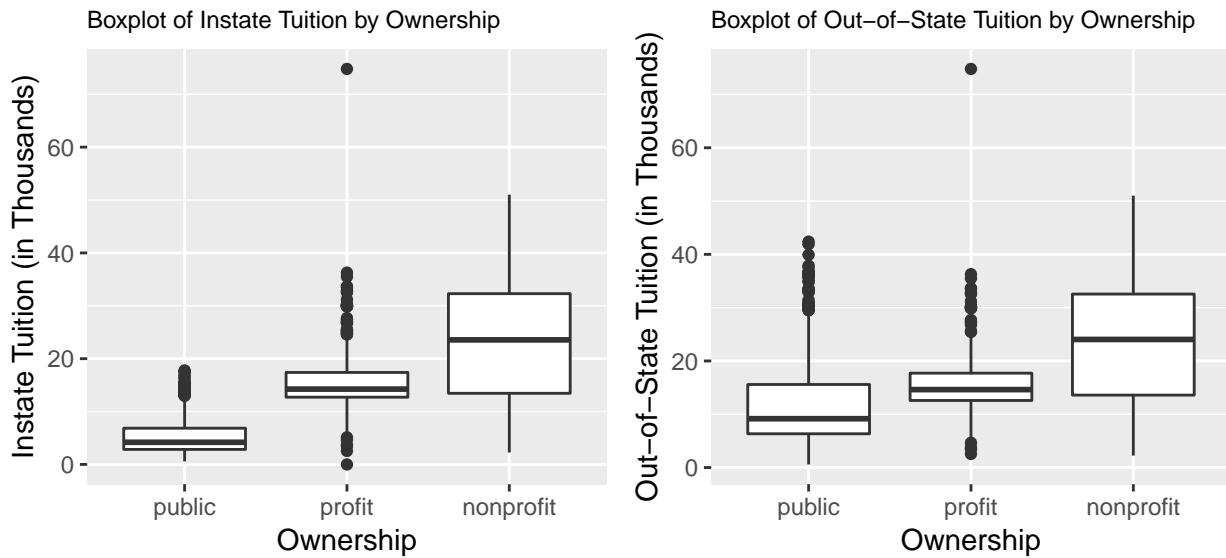


Histogram of Out-of-State Tuition by Ownership



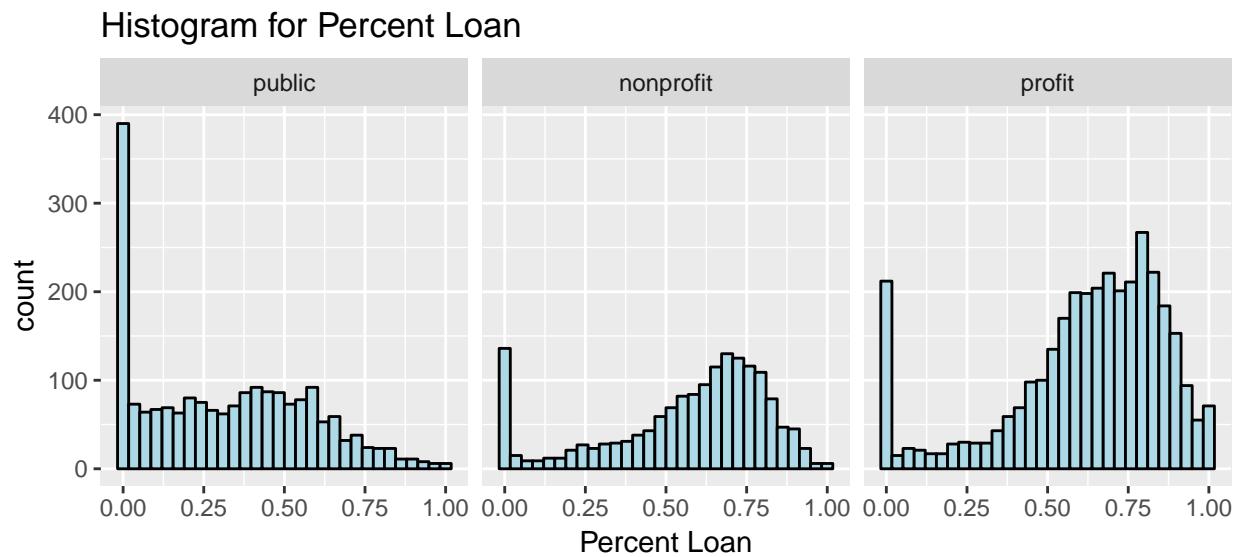
The boxplot for tuition by ownership below gives us a better understanding of the median, as well as outliers of the data. We can see that the median for private nonprofit schools is the highest, while that of public is the lowest. Moreover, we can clearly see the outliers. The one outlier for private for-private schools with value over 70,000 dollar seems like an unusual observation.

```
p7 <- ggplot(college1, aes(x=reorder(ownership,tuition_instate),y=tuition_instate))+  
  geom_boxplot() +  
  labs(y="Instate Tuition (in Thousands)",x="Ownership") +  
  ggtitle("Boxplot of Instate Tuition by Ownership") +  
  theme(plot.title = element_text(size = rel(0.8)))  
p8 <- ggplot(college2, aes(x=reorder(ownership,tuition_out),y=tuition_out))+  
  geom_boxplot() +  
  labs(y="Out-of-State Tuition (in Thousands)",x="Ownership") +  
  ggtitle("Boxplot of Out-of-State Tuition by Ownership") +  
  theme(plot.title = element_text(size = rel(0.8)))  
grid.arrange(p7,p8,ncol=2)
```



The percent loan variable represents percent of all undergraduate students receiving a federal student loan. By faceting into different ownership groups, the histogram shows that private for-profit schools have the highest rates of students receiving a federal loan, while public schools have the lowest. For all ownership types, there are a lot of students receiving no federal loan.

```
college3 = college[!is.na(college$pct_loan), ]
college3$ownership <- recode[college3$ownership]
college3$ownership<-fct_relevel(college3$ownership, 'public', 'nonprofit', 'profit')
ggplot(college3, aes(x=pct_loan))+
  geom_histogram(fill = "lightblue", color = "black")+
  labs(x="Percent Loan")+
  ggtitle("Histogram for Percent Loan")+
  facet_grid(.~ownership)
```



Values

Education Quality

More and more colleges are hiring part-time faculty as a means of cost-cutting. However, we should emphasize the importance of `full-time faculty` as it might be closely related to students' academic performance. For one thing, full-time faculty members are more stable in terms of curriculum teaching and planning. Because of high flexibility, they are also more likely to have higher levels of availability for students after class.

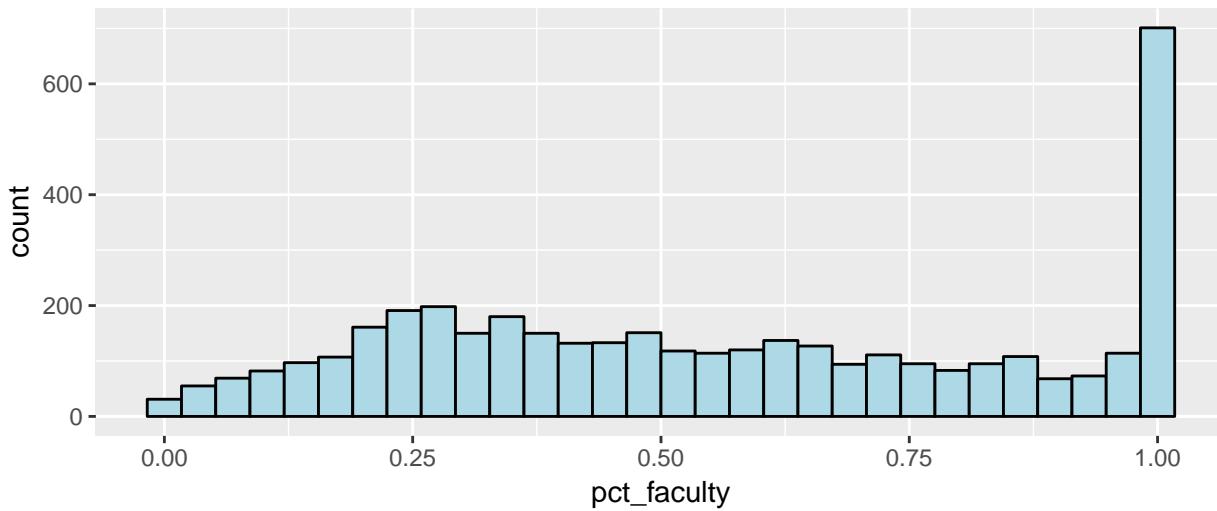
Using the feature `pct_faculty`, *percentage of full-time faculty*, we see that the distribution of percent of full-time faculty is slightly positively skewed with a mode of 1.00. Thus, most colleges have entirely full-time faculty members.

```
library(readr)
library(reshape)

college = read_rds("college.rds")
faculty = college %>% select(name, avg_faculty_salary, pct_faculty, num_undergrad, num_grad)

ggplot(faculty, aes(x = pct_faculty)) +
  geom_histogram(na.rm = TRUE, fill = "lightblue", color = "black")+
  labs(title = "Histogram for Full-time Faculty Percentage")
```

Histogram for Full-time Faculty Percentage

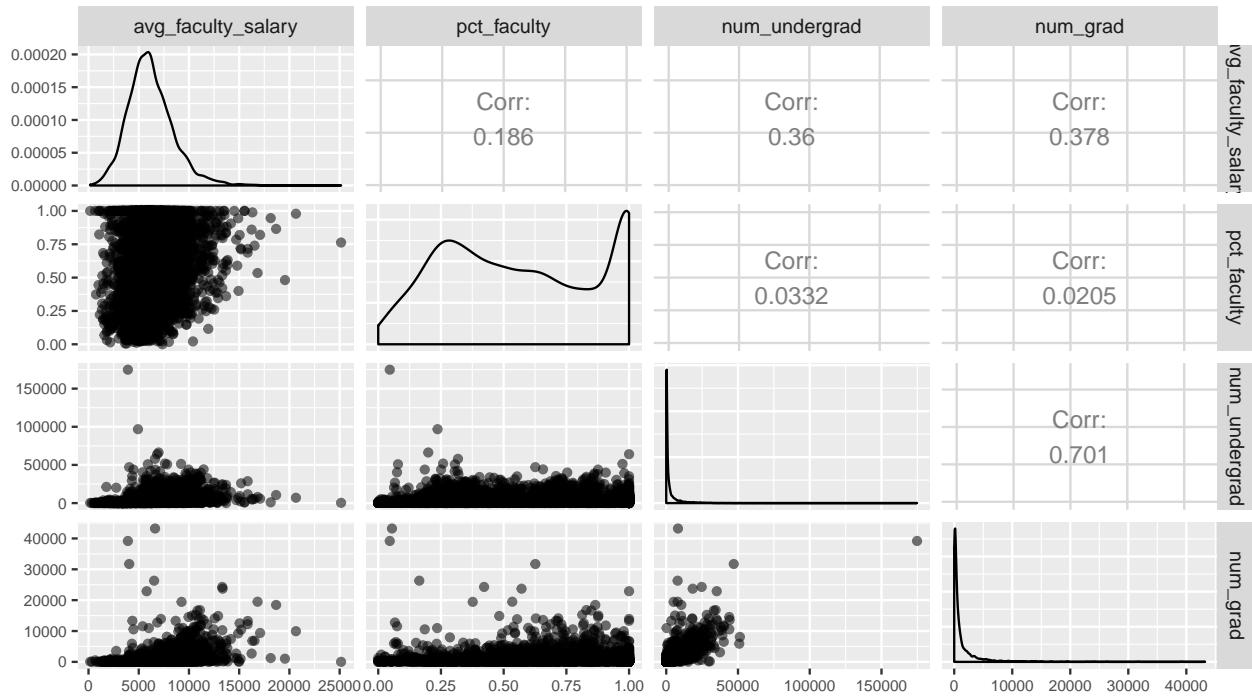


For further analysis of `full-time faculty`, we propose two research questions:

1. Is % of full-time faculty associated with the number of undergraduates/graduates?
2. Is % of full-time faculty associated with mean faculty salary?

To investigate the relationships, we obtained the following scatterplot matrix. We conclude that there is barely any relationship between percent of full-time faculty and other variables we proposed, as the correlations are negligible.

```
GGally::ggpairs(faculty, aes(alpha = 0.01), columns = colnames(faculty %>% select(-name))) +
  theme(axis.text.x = element_text(size = rel(0.8)),
        axis.text.y = element_text(size = rel(0.8)),
        strip.text = element_text(size = rel(0.8)))
```



Diversity

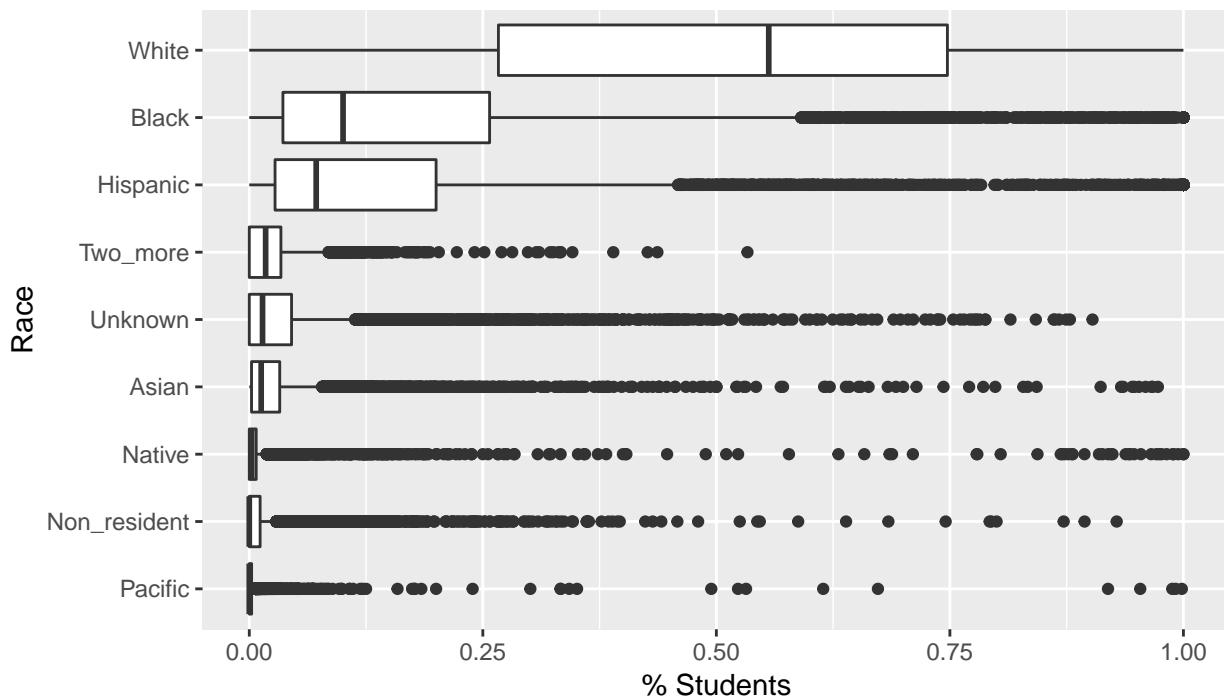
1. Racial Diversity

Racial segregation, with its root in the practice of slavery and segregation, is still present today in the United States. Racial separation impedes the establishment of a multicultural and equal society. In particular, it poses challenges for minority students to improve in educational settings, as discrimination, racial biases, and prejudices pervade. On the contrary, racial diversity can benefit students for it fosters collaboration, encourages creativity, and ultimately improves academic achievement. Therefore, we would like to provide an evaluation of racial diversity among colleges.

To begin the analysis, we first explore **how common each race is in colleges**.

```
college = read_rds("college.rds")
race_cat = c("race_white", "race_black", "race_hispanic",
            "race_asian", "race_native", "race_pacific",
            "race_2more", "race_nonresident", "race_unknown")
race = college %>% select ("name", race_cat)
colnames(race) = c("name", "White", "Black", "Hispanic", "Asian",
                  "Native", "Pacific", "Two_more", "Non_resident", "Unknown")
race_melted = melt(data.frame(race), id = "name")
g1 = ggplot(race_melted, aes(x = fct_reorder(`variable`, value, median, na.rm = TRUE),
                             y = value)) + geom_boxplot() + coord_flip() +
  labs(x = "Race", y = "% Students", title = "Percent of Students by Race")
g1
```

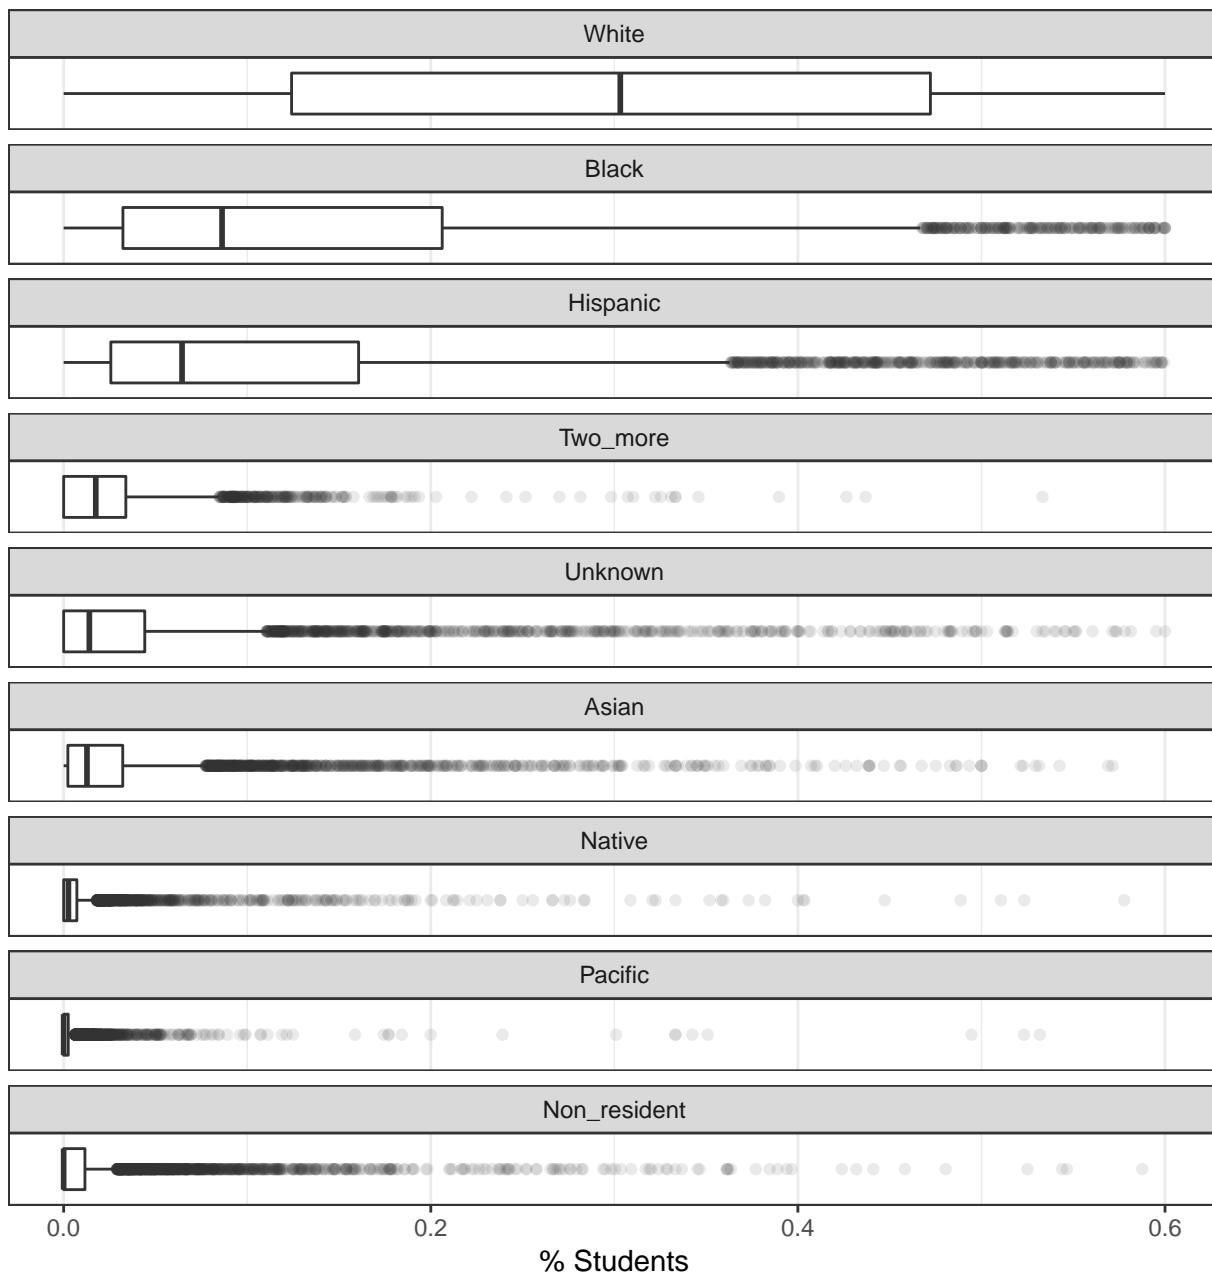
Percent of Students by Race



```
# Adjust for visualization
race_melted = race_melted %>%
  mutate(`variable` = fct_reorder(`variable`, value, median, na.rm = TRUE, .desc = TRUE))

# use "x = 1" to trick ggplot2 into making a single boxplot per facet
g2 = ggplot(race_melted, aes(x = 1, y = value)) +
  geom_boxplot(outlier.alpha = .1) +
  coord_flip() + facet_wrap(~`variable`, ncol = 1) +
  theme_bw() +
  scale_x_continuous(limits = c(0.4, 1.6)) +
  scale_y_continuous(limits = c(0, 0.6)) +
  # get rid of y-axis
  theme(axis.line.y = element_blank(),
        axis.ticks.y = element_blank(),
        axis.text.y = element_blank(),
        axis.title.y = element_blank(),
        panel.grid.major.y = element_blank(),
        panel.grid.minor.y = element_blank()) +
  labs(x = "Race", y = "% Students", title = "Percent of Students by Race (capped at 0.6)")
g2
```

Percent of Students by Race (capped at 0.6)



Comparing the medians in the boxplots we find the between-group variation to be significant. Specifically, **White**, **Black** and **Hispanic** seem to be the most common race while **Pacific** and **Non_resident** are the least. Also, in-group variation is larger for **White**, **Black** and **Hispanic** as well, so that the proportions of these racial groups vary more from college to college. Note that there are a great many outliers in the dataset.

Racial Diversity Index (RDI)

Next, to analyze racial diversity, we introduce *Racial Diversity Index (RDI)*, which is defined by the probability that any two students selected at random would have different races. This index was developed by Missouri

State University (<https://diversity.missouristate.edu/DiversityIndex.htm>) and is calculated by

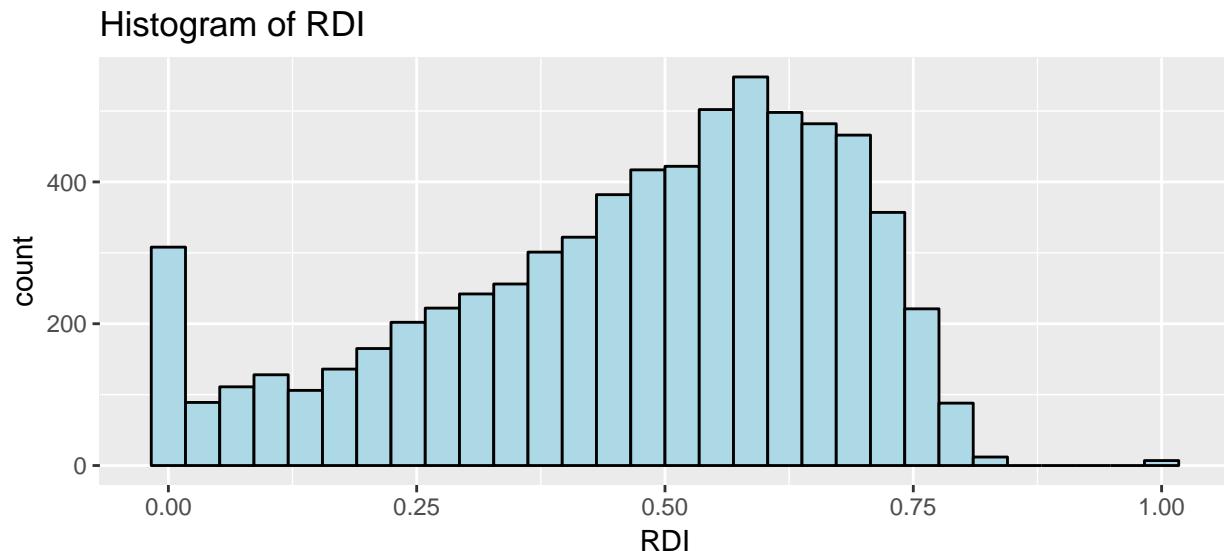
$$RDI = 1 - \Pr(\text{same race}) = 1 - \sum_i (\text{race } i)^2,$$

where *race i* is the percentage of students of *race i*.

```
race <- race %>% mutate(RDI = 1 - {(race$White)^2 + (race$Black)^2 +
  (race$Hispanic)^2 + (race$Asian)^2 + (race$Native)^2 + (race$Pacific)^2 +
  (race$Two_more)^2 + (race$Non_resident)^2 + (race$Unknown)^2})
```

```
g3 = ggplot(race, aes(x = RDI)) +
  geom_histogram(na.rm = TRUE, fill = "lightblue", color = "black") +
  labs(title = "Histogram of RDI")
```

```
g3
```



An RDI close to 1 implies high racial diversity while an RDI close to 0 implies low. From the histogram, we observe that the distribution of RDI is slightly negatively skewed, so that more colleges have moderate to high RDI's (over 0.50). The mode is approximately 0.6 so that most colleges have moderately high racial diversity.

Note that there are a few colleges scoring 0.00 (*Table 1*), as all of their students are from one exclusive race. For instance, all students of *Thunderbird School of Global Management* are *Hispanic*. At the other tail, note the isolated bar for RDI = 1.00 , which is impractical as no college would have exactly even distribution of races. A closer look at the data reveals that the *race* features for these colleges are all 0, indicating that their population might not have been recorded properly (*Table 2*). Excluding these colleges, we find the most racially diverse colleges listed in *Table 3* below.

```
library(knitr)
library(kableExtra)
kable(head(race %>% filter(RDI == 0)), format = "latex",
      caption = "Colleges with RDI = 0") %>%
  kable_styling(latex_options="scale_down")
```

```
kable(head(race %>% filter(RDI == 1)), format = "latex",
      caption = "Colleges with RDI = 1") %>%
  kable_styling(latex_options="scale_down")
```

Table 1: Colleges with RDI = 0

name	White	Black	Hispanic	Asian	Native	Pacific	Two_more	Non_resident	Unknown	RDI
Thunderbird School of Global Management	0	0	1	0	0	0	0	0	0	0
Velvatech College of Beauty Culture	0	1	0	0	0	0	0	0	0	0
Universal College of Beauty Inc-Los Angeles 1	0	1	0	0	0	0	0	0	0	0
Carver Bible College	0	1	0	0	0	0	0	0	0	0
Mr Leon's School of Hair Design-Moscow	1	0	0	0	0	0	0	0	0	0
Hebrew Theological College	1	0	0	0	0	0	0	0	0	0

Table 2: Colleges with RDI = 1

name	White	Black	Hispanic	Asian	Native	Pacific	Two_more	Non_resident	Unknown	RDI
American Conservatory Theater	0	0	0	0	0	0	0	0	0	1
Prince Institute-Rocky Mountains	0	0	0	0	0	0	0	0	0	1
Lyme Academy College of Fine Arts	0	0	0	0	0	0	0	0	0	1
Professional Business College	0	0	0	0	0	0	0	0	0	1
Spanish-American Institute	0	0	0	0	0	0	0	0	0	1
Taft University System	0	0	0	0	0	0	0	0	0	1

```
kable(head(race %>% filter(RDI != 1) %>% arrange(-RDI) %>% select(name, RDI)),
      format = "latex",
      caption = "Colleges with Highest Racial Diversity")
```

Finally, taking a look at the make-ups of the three most racially diverse colleges, we can see that the proportions of races can still vary a lot within each college. For instance, at *Chaminade University of Honolulu*, almost one-third of students are Asian but only few are Native and Non-resident.

```
top3 = c("Chaminade University of Honolulu", "Everest College-Renton",
        "Schiller International University")

race_top3 = race_melted %>% filter(name %in% top3)

g4 = ggplot(race_top3, aes(x = variable, y = value)) +
  geom_bar(stat="identity", position = position_dodge(),
            fill = "lightblue", color = "black") + facet_grid(.~ name) +
  labs(title = "Percent of Students by Race", x = "Race", y = "% Students") +
  theme(axis.ticks.x = element_blank(),
        axis.text.x = element_text(angle = 45, hjust = 1),
        panel.grid.major.x = element_blank(),
        text = element_text(size = 10)) +
  scale_y_continuous(breaks = seq(0, 0.4, 0.05))

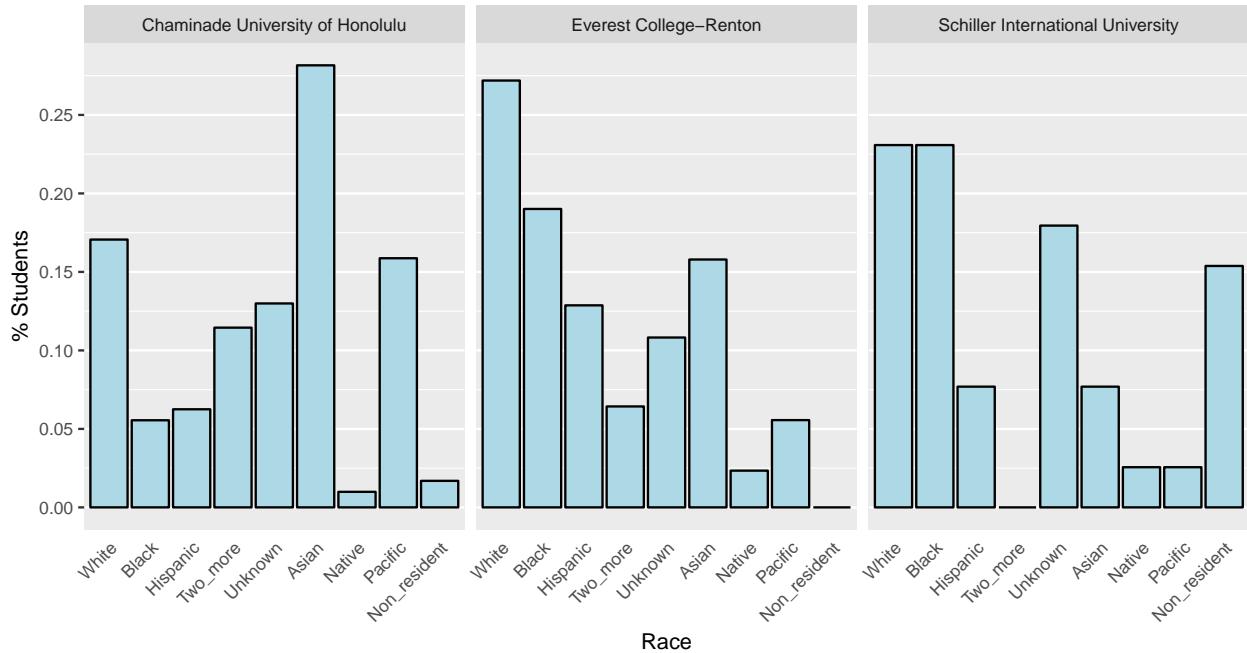
#Remark. Some questions for further exploration: is *racial diversity* correlated to other variables such as `ownership`, `predominant_Degree` and `highest_degree`?
```

Table 3: Colleges with Highest Racial Diversity

name	RDI
Chaminade University of Honolulu	0.8290570
Everest College-Renton	0.8289556
Schiller International University	0.8244501
Pacific Rim Christian University	0.8215979
Golden Gate University-San Francisco	0.8205131
Charles A Jones Career and Education Center	0.8203344

g4

Percent of Students by Race

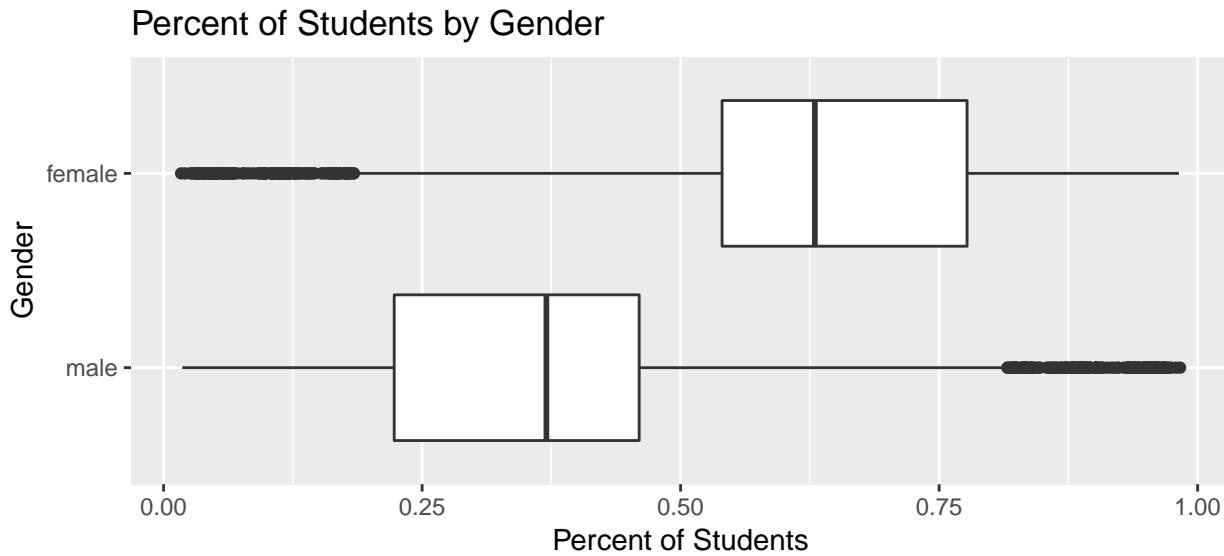


2. Gender Diversity

Gender equality has been an issue in the workplace, where women are known to have struggled for parity in pay and influence. To form awareness of gender equality and reduce challenges in workplace in the future, more and more colleges have been working on promoting gender diversity.

To investigate gender diversity, we first explore **how common each gender is in colleges**. Comparing the medians in the boxplot we see that most colleges seem to have more female students than male.

```
gender = college %>% select(name, pct_female) %>%
  mutate(pct_male = 1-pct_female)
colnames(gender) = c("name", "female", "male")
gender_melted = melt(data.frame(gender), id = "name")
g5 = ggplot(gender_melted, aes(x = fct_reorder(`variable`, value, median, na.rm = TRUE),
                                y = value)) + geom_boxplot() + coord_flip() +
  labs(x = "Gender", y = "Percent of Students", title = "Percent of Students by Gender")
g5
```



Gender Diversity Index (GDI)

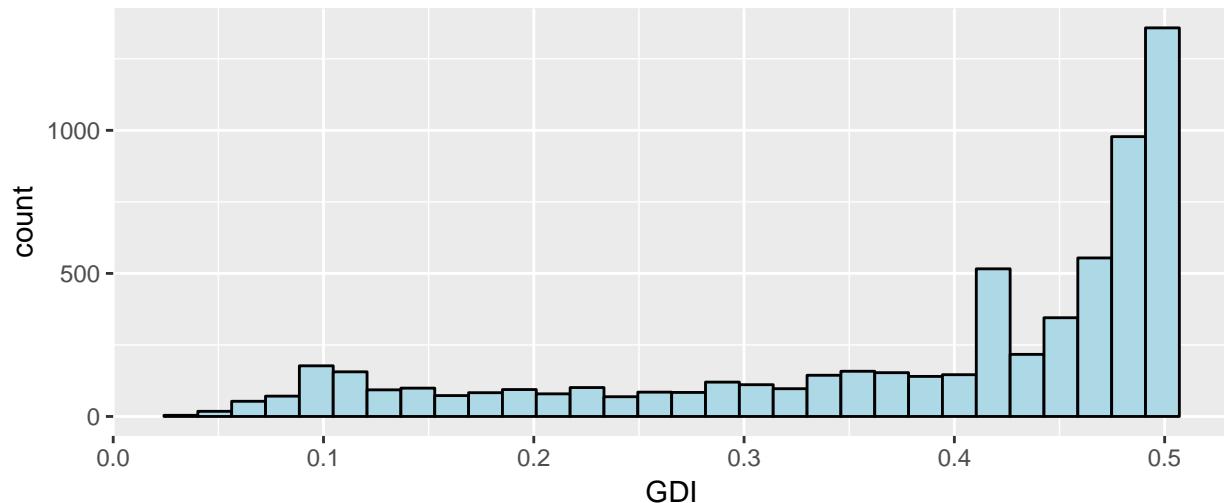
To give a substantial measurement of gender diversity, we develop the *Gender Diversity Index (GDI)*. Like **racial diversity**, it is defined by the probability that any two students selected at random would have different genders:

$$GDI = 1 - \Pr(\text{same gender}) = 1 - (\text{male})^2 - (\text{female})^2,$$

where *male* and *female* are percent of male and female students.

```
gender = gender %>% mutate(GDI = 1-female^2-male^2)
ggplot(gender, aes(x = GDI)) +
  geom_histogram(na.rm = TRUE, fill = "lightblue", color = "black")+
  labs(title = "Histogram of GDI")
```

Histogram of GDI



A GDI close to 0.5 implies high racial diversity (i.e., the percentage of male students are close to percentage of female students) while a GDI close to 0 implies low racial diversity. From the histogram, we observe that the distribution of GDI is negatively skewed with a long left tail. Note that the mode is 0.5, so that most colleges have equal % of female and male students.

Table 4: Colleges with Highest Gender Diversity

name	female	male	GDI
VanderCook College of Music	0.5000000	0.5000000	0.5
South Carolina State University	0.5000000	0.5000000	0.5
American Institute of Beauty	0.5000000	0.5000000	0.5
SOLEX College	0.5000000	0.5000000	0.5
Texas State Technical College-Harlingen	0.5001274	0.4998726	0.5
University of Idaho	0.5001372	0.4998628	0.5

Finally, the colleges with the highest gender diversity can be found in Table 4.

```
kable(head(na.omit(gender)) %>% arrange(-GDI)),
  caption = "Colleges with Highest Gender Diversity")
```

Completion

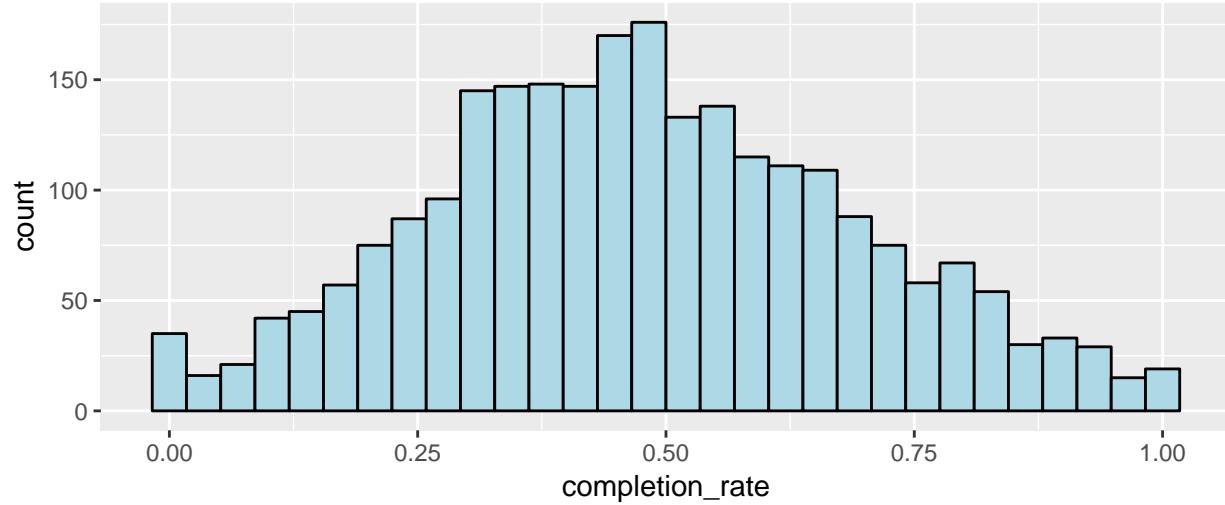
1. Completion Rate

Completion rate is recorded as *percentage of students that graduated within 6 years* in our dataset. The overall distribution of completion rate is roughly symmetric, with mode around 0.47. Thus, for most colleges, just under half (47%) of students graduated within 6 years of starting college.

```
college = read_rds("college.rds")

ggplot(college, aes(x = completion_rate)) +
  geom_histogram(na.rm = TRUE, fill = "lightblue", color = "black")+
  labs(title = "Histogram of Completion Rate")
```

Histogram of Completion Rate



Completion Rate by Institution Type

Taking institution type into consideration, we can see that private *non-profit* schools have the best completion rates (highest median) while private *for-profit* schools have the worst. Only non-profit schools have higher median completion rate than the overall median. Note that unlike *non-profit* and *for-profit* schools, public schools have neither a completion rate of 0.00 or 1.00.

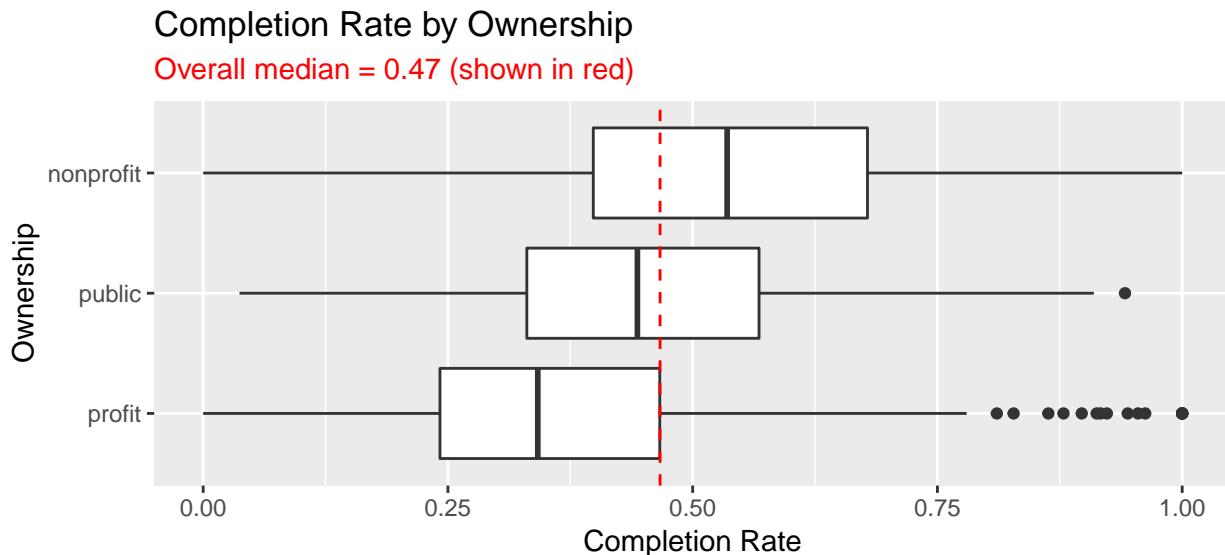
```

comp_rate = college %>% select(name, ownership, completion_rate)
levels(comp_rate$ownership) = c("public", "nonprofit", "profit")

med = median(comp_rate$completion_rate, na.rm = TRUE)

ggplot(comp_rate, aes(x = fct_reorder(`ownership`, completion_rate, median, na.rm = TRUE),
                      y = completion_rate)) + geom_boxplot() + coord_flip() +
  geom_hline(yintercept = med, color = "red", linetype = 2) +
  labs(x = "Ownership", y = "Completion Rate", title = "Completion Rate by Ownership",
       subtitle = paste0("Overall median = ", round(med, 2), " (shown in red)")) +
  theme(plot.subtitle = element_text(color = "red"))

```



2. Degree Awarded

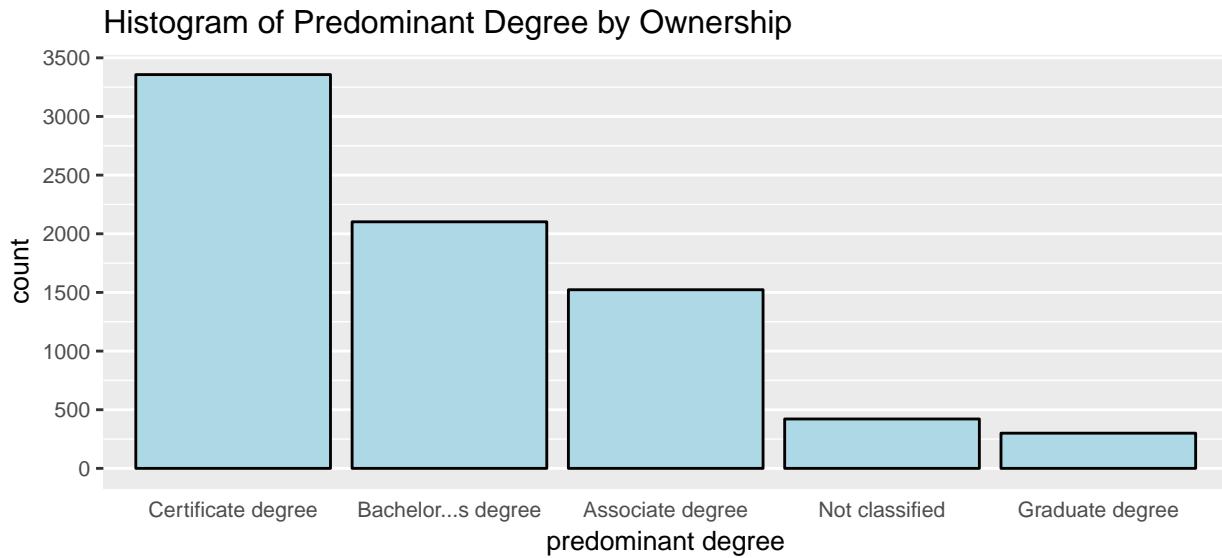
Overall, *Certificate degree* is the most frequently awarded while *Graduate degree* is the least. Such a result is reasonable as certificate programs are less demanding than degree programs. Also, since graduate population is usually smaller than undergraduate population, we shall expect that fewer students receive a *Graduate degree*.

```

degree = college %>% select(name, ownership, predominant_degree, highest_degree)
levels(degree$ownership) = c("public", "nonprofit", "profit")
levels(degree$predominant_degree) = c("Not classified", "Certificate degree", "Associate degree", "Bachelor's degree", "Graduate degree")
levels(degree$highest_degree) = c("Non-degree-granting", "Certificate degree", "Associate degree", "Bachelor's degree", "Graduate degree")

gc3 = ggplot(degree, aes(x = fct_reorder(predominant_degree, predominant_degree, function(x){length(x)}))
              geom_bar(fill = "lightblue", color = "black") +
              labs(x = "predominant degree") +
              theme(axis.ticks.x = element_blank(),
                    panel.grid.major.x = element_blank(),
                    text = element_text(size = 10)) +
              scale_y_continuous(breaks = seq(0, 3500, 500))+
              labs(title = "Histogram of Predominant Degree by Ownership")
gc3

```



Predominant Degree by Ownership

Predominant degree seems to be associated with ownership as the boxes in the mosaic plot below are not proportional. Specifically, we can infer that the predominant degree awarded is mostly

- *Certificate degree* for private *for-profit* schools;
- *Associate degree* for *public* schools;
- *Bachelor's degree* for private *non-profit* schools.

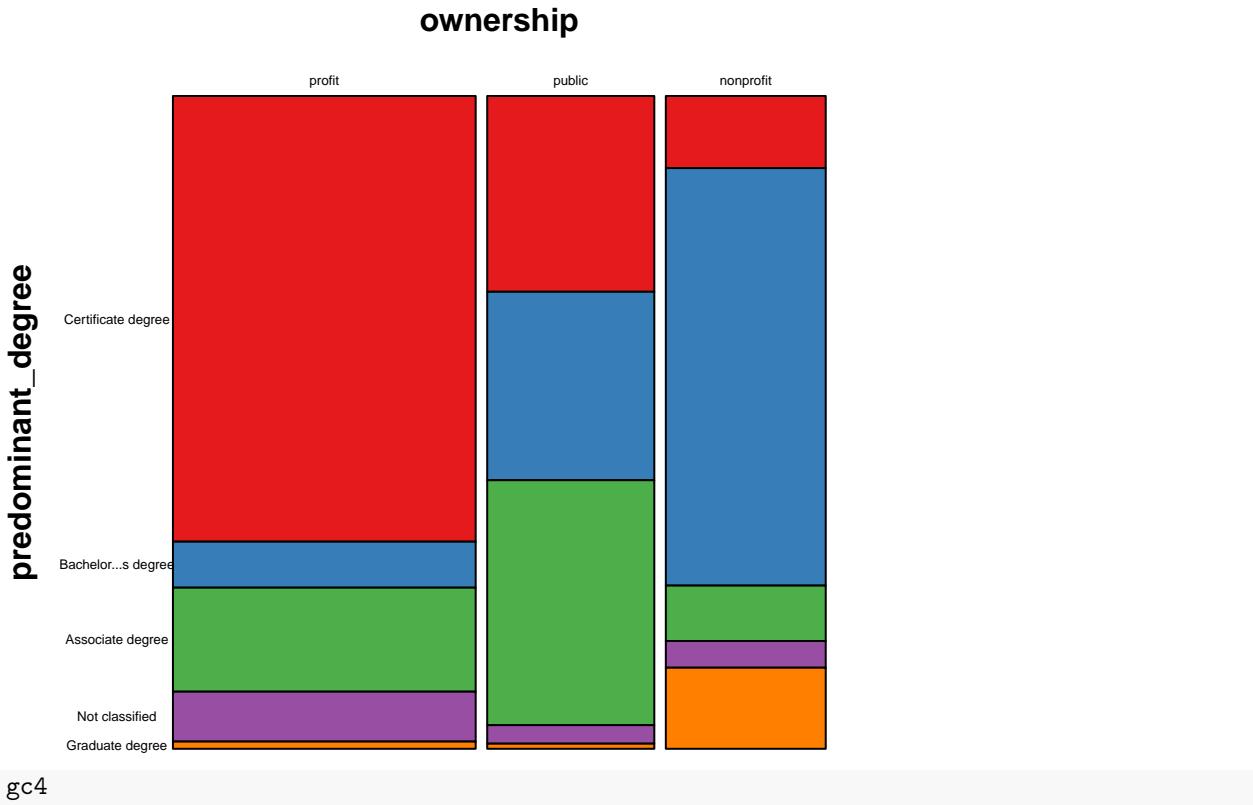
```
library(grid)
library(vcd)
library(RColorBrewer)
fillcolors = brewer.pal(5, "Set1")

# Sort predominant_degree by frequency
count1 = degree %>% group_by(ownership, predominant_degree) %>% summarize(Freq = n())
pred_deg_order = count1 %>% group_by(predominant_degree) %>%
  summarize(sum = sum(Freq)) %>% arrange(-sum) %>%
  pull(predominant_degree)
count1$predominant_degree = factor(count1$predominant_degree, levels = pred_deg_order)

# Sort ownership by frequency
ownership_order = count1 %>% group_by(ownership) %>%
  summarize(sum = sum(Freq)) %>% arrange(-sum) %>%
  pull(ownership)
count1$ownership = factor(count1$ownership, levels = ownership_order)

gc4 = mosaic(predominant_degree~ownership, count1, direction = c("v", "h"),
             gp = gpar(fill = fillcolors),
             labeling = labeling_border(gp_labels = gpar(fontsize = 5),
                                       rot_labels = c(0, 0, 0, 0),
                                       offset_varnames = c(0,0,0,2),
                                       offset_labels = c(0,0,0,2.5)),
             main = "Predominant Degree vs. Ownership")
```

Predominant Degree vs. Ownership



gc4

```
##                                     ownership profit public nonprofit
## predominant_degree
## Certificate degree            2528     613      216
## Bachelor's degree             261      590     1251
## Associate degree              590      767      166
## Not classified                283      58       80
## Graduate degree                 41      16      243
```

Earnings

Although earning is not the only criteria to see how successful a person is, it is a very straightforward indication of performance after graduation. Thus, we examine earnings after graduation as a part of values a college provides.

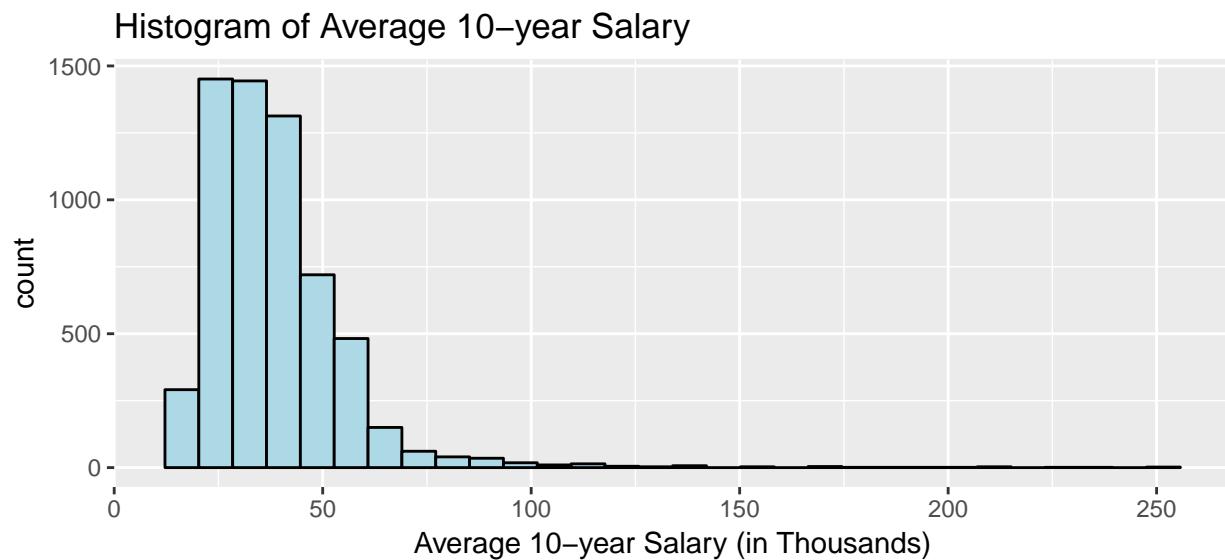
The `average 10-year salary` variable represents the mean earnings of students working and not enrolled at school 10 years after entry. The overall distribution of the average 10 year salary is skewed to the right. Most of our observations have an average 10-year salary of around 20,000 to 40,000 dollars. There are also a couple of observations with higher average 10 year salary than 100,000 dollars.

```
college = read_rds("college.rds")
college[college=='NULL'] <- NA
#college = na.omit(college$avg_10yr_salary)
college = college[!is.na(college$avg_10yr_salary), ]
college$avg_10yr_salary = college$avg_10yr_salary/1000
college$tuition_instate = college$tuition_instate/1000
college$tuition_out = college$tuition_out/1000
```

```

recode <- c('1'='public', '2'='nonprofit', '3'='profit')
college$ownership <- recode[college$ownership]
college$ownership<-fct_relevel(college$ownership, 'public','nonprofit','profit')
ggplot(college, aes(x=avg_10yr_salary))+
  geom_histogram(bins = 30,fill = "lightblue", color = "black")+
  labs(x="Average 10-year Salary (in Thousands)")+
  ggtitle("Histogram of Average 10-year Salary")

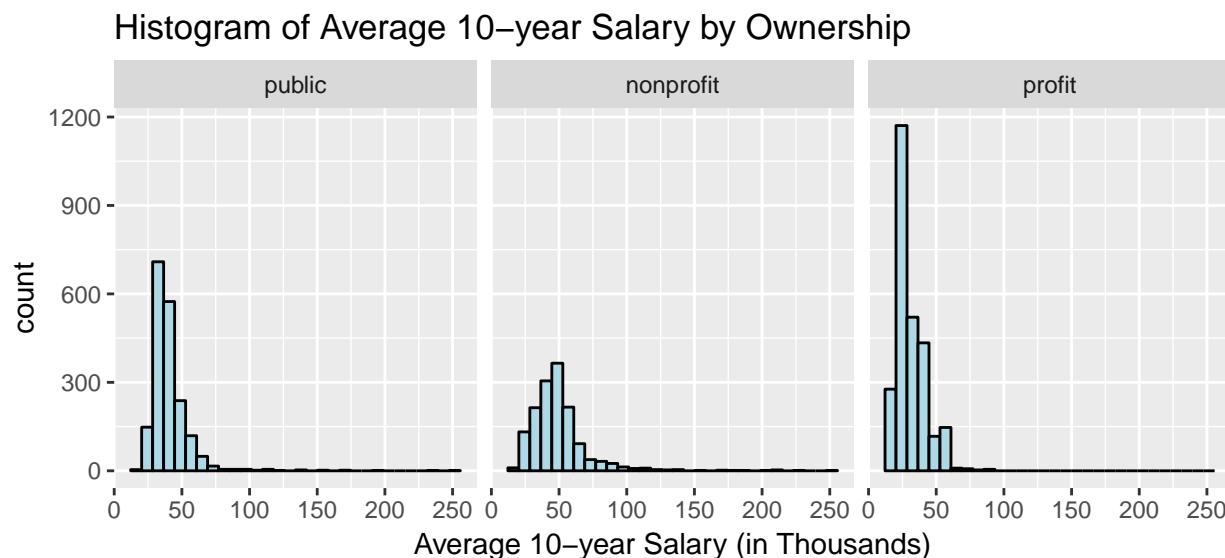
```



```

a=count(college[ which(college$ownership=="public"),])
b=count(college[ which(college$ownership=="nonprofit"),])
c=count(college[ which(college$ownership=="profit"),])
ggplot(college, aes(x=avg_10yr_salary))+
  geom_histogram(bins = 30,fill = "lightblue", color = "black")+
  facet_grid(.~ownership)+
  labs(x="Average 10-year Salary (in Thousands)")+
  ggtitle("Histogram of Average 10-year Salary by Ownership")

```

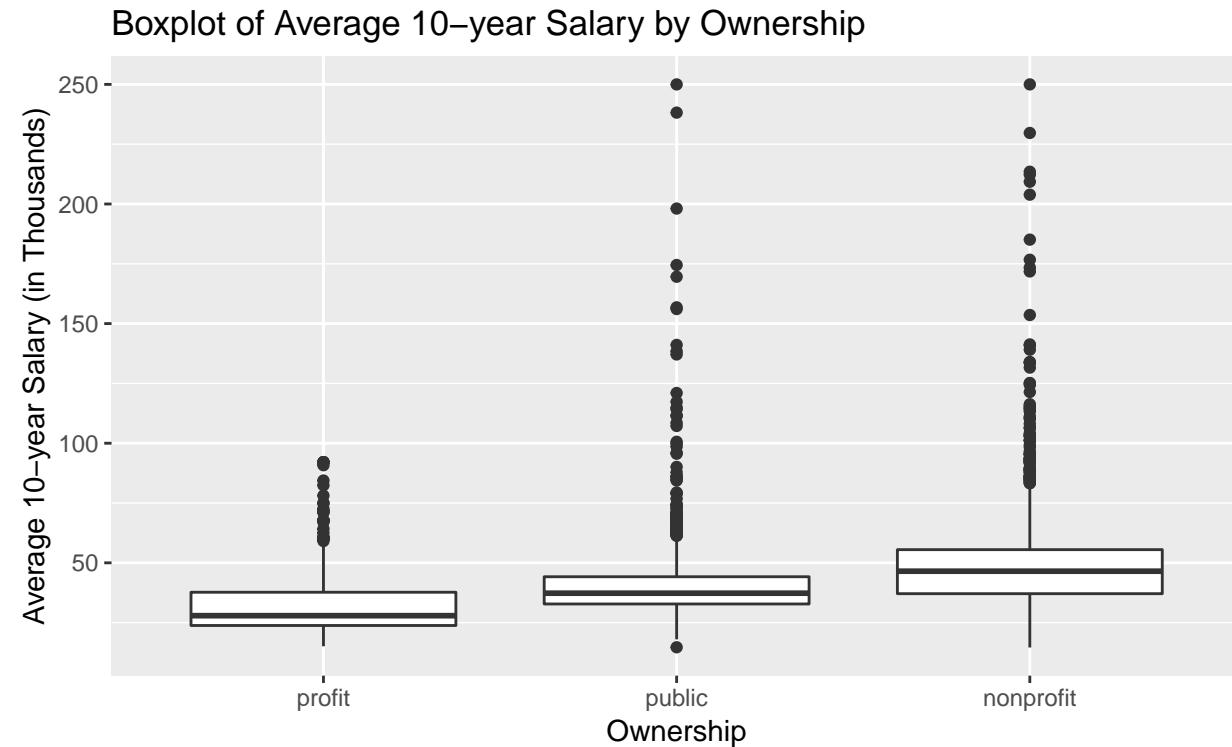


By faceting average 10-year salary into different ownership groups, we can see the different distributions

for public, private nonprofit, and private for-profit schools. The total numbers of observations for the three ownership type are 1890, 1481, 2691. Among the three types, although private non-profit schools have the least observations, they have the distribution the most skewed to the right, with the highest mean and mode. This indicates students from private non-profit schools have highest average 10 year salary among the three types of schools. While public schools have the mode for average 10-year salary at around 30,000 dollars, private for-profit schools have the mode at around 20,000 dollars, and overall, students from public schools rank the second in average 10 year salary. Therefore, the average 10-year salary differs for students from schools with different ownership types.

The boxplot below gives us a better understanding of the median, as well as outliers in the data. We can see that the median for private non-profit schools is the highest, while that of private for-profit ones is the lowest. Moreover, we can clearly see the outliers. Comparing to other ownership types, private for-profit schools have the least number of outliers, and the value of outlier is the lowest (no outliers with value over 100,000). The boxplot also indicates students from private non-profit schools have higher average 10-year salary than those from public and private for-profit schools.

```
ggplot(college, aes(x=reorder(ownership,avg_10yr_salary),y=avg_10yr_salary))+
  geom_boxplot()+
  labs(y="Average 10-year Salary (in Thousands)",x="Ownership")+
  ggtitle("Boxplot of Average 10-year Salary by Ownership")
```



Cost-Value Tradeoffs

Eventually, to explore the cost-value trade-off, we obtained the following scatterplot matrix. We conclude the correlation to be *strong*, if $r \geq 0.7$; *weak*, if $r \leq 0.3$; *moderate* otherwise. Some key observations are:

- **Multicollinearity** exists for both the cost and value variables such that
 - `sat_avg` is moderately positively associated with `tuition_instate/tuition_out`
 - `sat_avg` is moderately negatively associated with `admission_rate`
 - `tuition_instate` is strongly positively associated with `tuition_out`
 - `completion_rate` is moderately positively associated with `avg_10yr_salary`

- `sat_avg` is strongly positively associated with both `completion_rate` and `avg_10yr_salary`
- `admission_rate` is moderately negatively associated with both `completion_rate` and `avg_10yr_salary`
- `tuition_instate` is moderately positively associated with both `completion_rate` and `avg_10yr_salary`
- `tuition_out` is moderately positively associated with both `completion_rate` and `avg_10yr_salary`

```

library(GGally)
college = read_rds("college.rds")

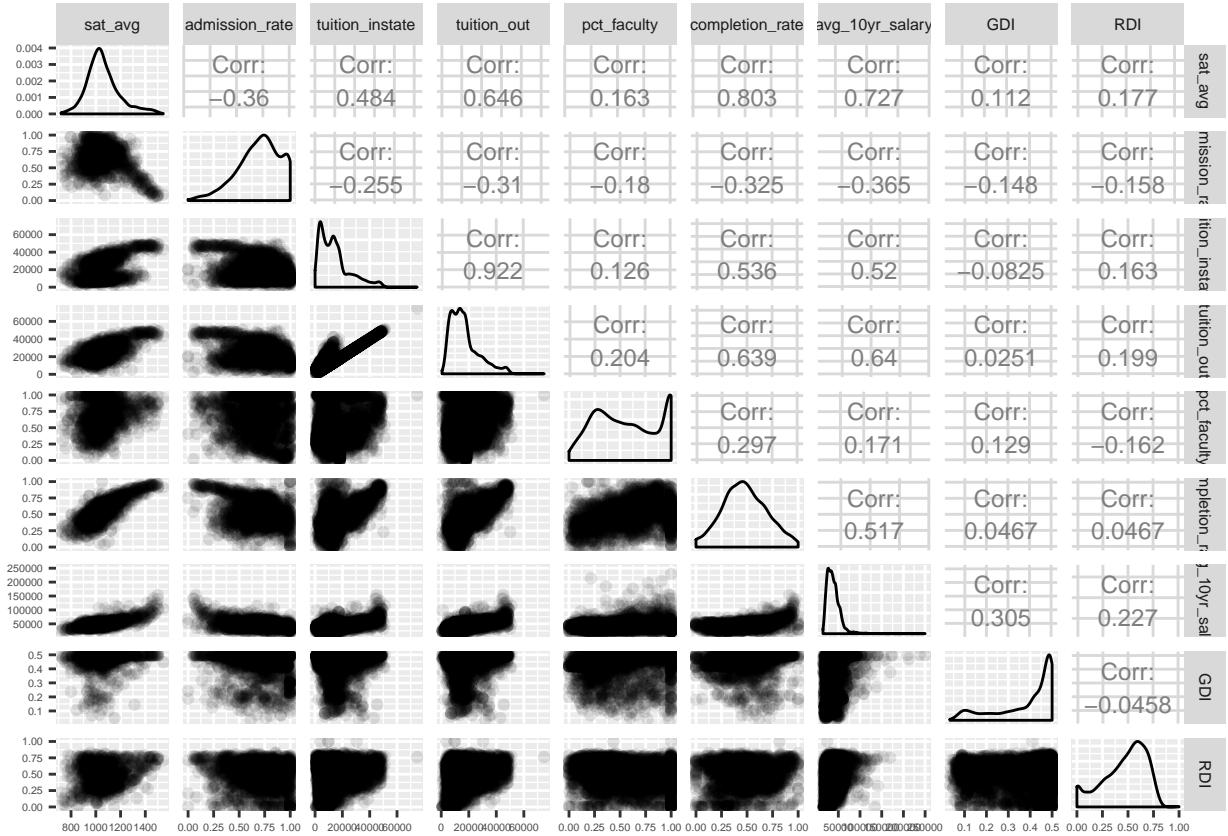
##### For individual file: duplicate code with Diversity.Rmd
race_cat = c("race_white", "race_black", "race_hispanic",
            "race_asian", "race_native", "race_pacific",
            "race_2more", "race_nonresident", "race_unknown")
race = college %>% select ("name", race_cat)
colnames(race) = c("name", "White", "Black", "Hispanic", "Asian",
                  "Native", "Pacific", "Two_more", "Non_resident", "Unknown")
race = race %>% mutate(RDI = 1 - {(race$White)^2 + (race$Black)^2 +
  (race$Hispanic)^2 + (race$Asian)^2 + (race$Native)^2 + (race$Pacific)^2 +
  (race$Two_more)^2 + (race$Non_resident)^2 + (race$Unknown)^2})
gender = college %>% select(name, pct_female) %>%
  mutate(pct_male = 1-pct_female)
colnames(gender) = c("name", "female", "male")
gender = gender %>% mutate(GDI = 1-female^2-male^2)
#####

cost_value = college %>%
  select("sat_avg", "admission_rate", "tuition_instate",
         "tuition_out", "pct_faculty", "completion_rate", "avg_10yr_salary") %>%
  mutate(GDI = gender$GDI) %>%
  mutate(RDI = race$RDI)

cost = c("sat_avg", "admission_rate", "tuition_instate", "tuition_out")
value = c("pct_faculty", "completion_rate", "avg_10yr_salary", "GDI", "RDI")

ggpairs(cost_value,
        upper = list(continuous = wrap("cor", size = 3, alignPercent = 1)),
        lower = list(continuous = wrap("points", alpha = 0.1))) +
  theme(axis.text.x = element_text(size = rel(0.5)),
        axis.text.y = element_text(size = rel(0.5)),
        strip.text = element_text(size = rel(0.5)))

```

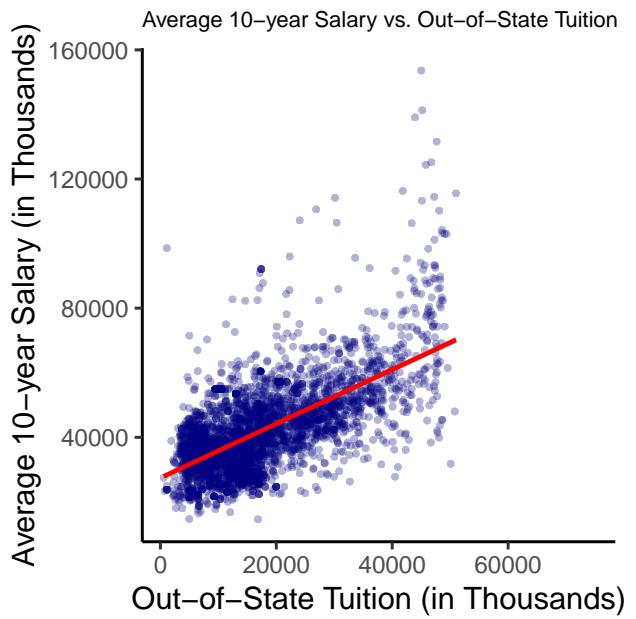
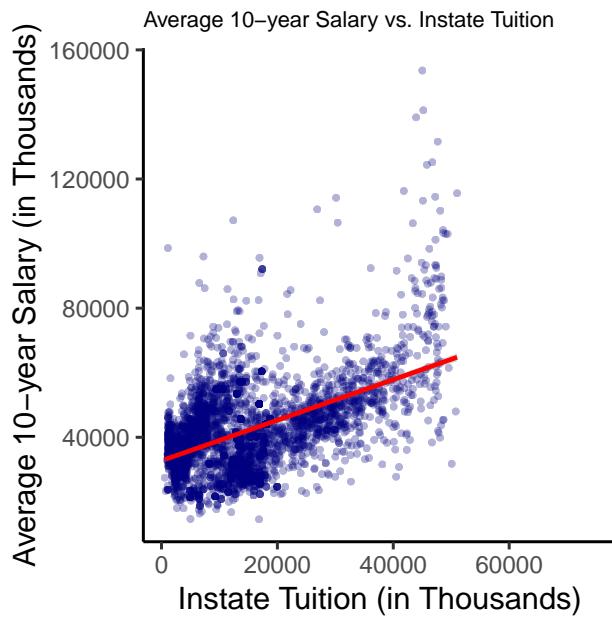


For our interest, we investigate the correlation between `tuition` and `average 10-year salary` in more details. We can see there is a positive relationship between `tuition` and `average 10-year salary`.

```

college = college[!is.na(college$tuition_out), ]
p1 <- ggplot(college, aes(x=tuition_instate,y=avg_10yr_salary))+
  geom_point(alpha = .3, color = "navy", stroke = 0)+
  labs(x="Instate Tuition (in Thousands)",y="Average 10-year Salary (in Thousands)")+
  ggtitle("Average 10-year Salary vs. Instate Tuition")+
  geom_smooth(method='lm',formula=y~x,color="red")+
  theme_classic(14)+
  theme(plot.title = element_text(size = rel(0.7)))
p2 <- ggplot(college, aes(x=tuition_out,y=avg_10yr_salary))+
  geom_point(alpha = .3, color = "navy", stroke = 0)+
  labs(x="Out-of-State Tuition (in Thousands)",y="Average 10-year Salary (in Thousands)")+
  ggtitle("Average 10-year Salary vs. Out-of-State Tuition")+
  geom_smooth(method='lm',formula=y~x,color="red")+
  theme_classic(14)+
  theme(plot.title = element_text(size = rel(0.7)))
r2_1 = format(summary(lm(college$avg_10yr_salary~college$tuition_instate))$r.squared*100, digits = 3)
r2_2 = format(summary(lm(college$avg_10yr_salary~college$tuition_out))$r.squared*100, digits = 3)
grid.arrange(p1,p2,ncol=2)

```



Additionally, we find that 28 and 40.9 percents of the variation in average 10-year salary are explained by the variation in instate tuition and out-of-state tuition, respectively.

EXECUTIVE SUMMARY

Our project provides analysis of different colleges to help high school students and their family to make informed decisions for college selection based on **cost-value** tradeoffs, where **cost** includes (1) *admission requirement* and (2) *affordability*, and **value** includes (1) *education quality*, (2) *diversity*, (3) *completion* and (4) *earnings*. The major findings for each factor and the relationship between costs and values will be discussed below.

Admission requirement. Admission requirement is assessed by **SAT score** and **admission rate**. Importantly, for schools with average SAT below 1200, admission rate doesn't seem to have strong relationship with average SAT scores. On the other hand, for those with average SAT higher than 1200, admission rate appears to decrease with higher SAT score. Such academically demanding schools may have an explicit requirement for SAT score for admission.

Affordability. Affordability is analyzed through **tuition** (out-of-state/instate) and **federal loan**. Generally, while out-of-state tuition is higher than instate tuition, schools with higher out-of-state tuition tend to have higher instate tuition. The distributions for out-of-state and instate tuition are similar across ownership groups. Public schools have the lowest tuition whereas private non-profit schools have the highest with larger variation. For federal loan, while private for-profit schools have the highest percentage of all undergraduate students receiving a federal loan, public schools have lowest. Therefore, we recommend that high schoolers with a limited budget to attend public schools where tuition is comparatively low, and private for-profit schools where they are more likely to be granted a federal loan.

Education quality. **Full-time faculty percentage** is selected as the indicator for education quality since it is closely related to students' academic performances. Most colleges have 100% full-time faculty while the rest have varying percentage from none to full.

Diversity. **Racial and gender diversity** are examined by Racial Diversity Index (RDI) and Gender Diversity Index (GDI), each of which represents the probability that any two students selected at random would have different races or genders. Colleges with the highest RDI's and GDI's are identified in the previous report.

Completion. Completion is evaluated by **completion rate** and **predominant degree awarded**. For most colleges, completion rate is not satisfactory with just under half of the students graduating within 6 years of starting colleges. The most frequently awarded predominant degree is Certificate, while the least is Graduate. Additionally, ownership seems to affect completion. We find that private non-profit schools have the highest completion rate, while private for-profit ones have the lowest. The predominant degree awarded are mostly Certificate, Bachelor's and Associate degree, for private for-profit, private non-profit and public schools, respectively. Therefore, depending on the type of degree one wants to pursue, one can choose colleges of different ownership types accordingly.

Earnings. We investigate earnings by **average 10-year salary** after entry under different ownership types. Private non-profit schools have the highest median for average 10-year salary, while private for-profit schools have the lowest. Additionally, there seems to be a positive relationship between tuition and average 10-year salary, so that students who paid higher tuition at college had higher average 10-year salary.

Cost-value trade-off. Lastly, we inspect cost-value trade-off using the scatterplot matrix. Most importantly, higher **average SAT score** and **tuition** seem to imply higher **completion rate** and **average 10-year salary**. Thus, we have evidence to state that schools with **higher cost** to attend will lead to **higher value**. The major cost-value trade-off is visualized below.

```
college = read_rds("college.rds")

cost_value = college %>%
  select("sat_avg", "tuition_instate", "tuition_out", "completion_rate", "avg_10yr_salary")

ggpairs(cost_value,
       upper = list(continuous = wrap("cor", size = 5, alignPercent = 1)),
```

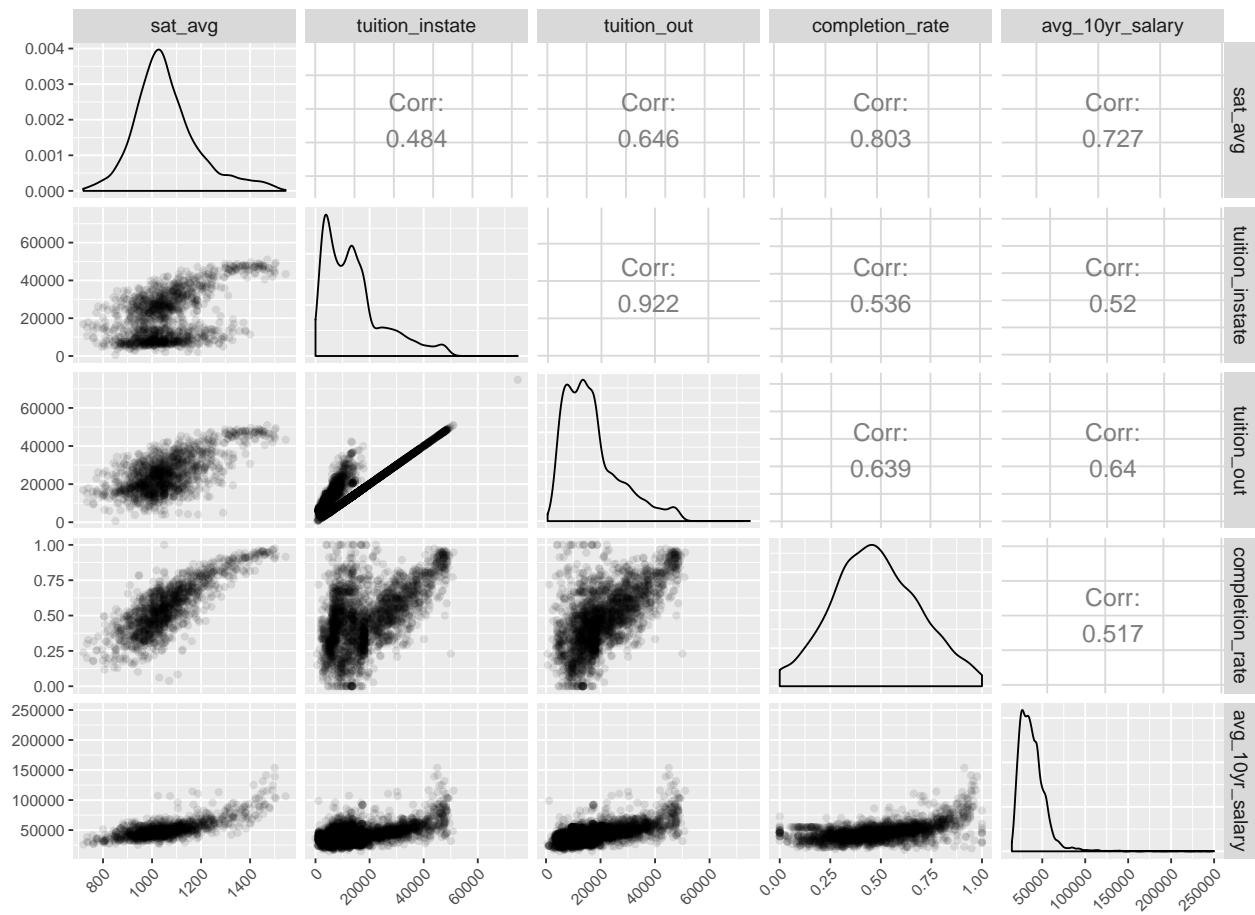
```

        lower = list(continuous = wrap("points", alpha = 0.1))) +
theme(axis.text.x = element_text(size = rel(1), angle = 45, hjust = 1),
      axis.text.y = element_text(size = rel(1)),
      strip.text = element_text(size = rel(1)),
      plot.title = element_text(size=18)) +
labs(title = "Scatterplot Matrix for Cost-Value Trade-offs",
     subtitle="Cost: Average SAT Score, Instate Tuition, Out-of-State Tuition
Value: Completion Rate, Average 10-year Salary")

```

Scatterplot Matrix for Cost–Value Trade–offs

Cost: Average SAT Score, Instate Tuition, Out-of–State Tuition
Value: Completion Rate, Average 10-year Salary



INTERACTIVE COMPONENT

Our interactive dashboard (https://yourbestcollegechoice.shinyapps.io/final_proj/) aims to help users to filter the characteristics of a specific university they are interested in and to make comparisons across universities. For example, if a person is interested in Columbia University, he can select Columbia University in the dashboard and see where Columbia University falls on each of the distribution plots of variables including education quality, diversity, completion rate and earnings. From the plots, users can easily compare Columbia University with other universities.

We choose Shiny as our interactive plot tool based on the following reasons. Firstly, Shiny is an effective high-level tool for the interactive plot. Shiny provides very convenient functions for creating a dashboard including base components like sidebar, main body, layout, and reactive plot from data change. Secondly, since Shiny is an R-based GUI programming tool, we could base our shiny app on the R codes we already have. Instead of spending much time on the grammar of a new language, we could focus on the logical part and user experience of this interactive dashboard.

A simple introduction to our Shiny Dashboard is as follows.

1. User cases

- Our research target is to find how schools with higher admission requirement and higher tuition fees correlate to higher education quality, diversity, completion rate, and earnings. We already have researched each variable and their relationships respectively. That means, now for a student or his/her parent, he has a sense of how these variables concerning with cost and value are distributed. However, he will ultimately focus on several universities that he can or wants to go to. Therefore, we provide an interactive tool for students and parents to narrow down the schools he will probably go to and see clearly the aspects of education quality, diversity, completion rate and earnings of that specific university.

2. Functions

- *Search universities by SAT scores and admission rate.* SAT is a requirement on admission to most universities and admission rate is an essential aspect students would consider when applying to a university. Therefore, we let users select the universities in some SAT score and admission rate range on the scatterplot, so that they could narrow down the range of universities they are able to attend.
- *Search a university by name.* Users could type the name of the university in the university table to access the info of that university directly.
- *Plot distribution on education quality, diversity, completion rate and earnings.* When a selection of a specific university has been made, the distribution on education quality, diversity, completion rate and earnings will be shown in the form of histograms. Where that university falls in the histogram is also visualized for a quick search.
- *Data exploring.* We provide a separate page for the data we use, because we want to enable users to explore the original data themselves. We also provide the metadata information and data missing patterns for reference.

3. User instructions

- Draw a square range based on the SAT scores and admission rate on the scatterplot. A list of universities within that SAT range and admission rate range will be shown. Then click a university that you are interested in in the table. The distribution plots with the specified university highlighted will be shown below. Besides, you can go to the data page to explore the metadata information, the original data and the data missing patterns.

4. Code links

- https://github.com/ZhangZhida/EDAV_FinalProject/blob/master/shiny/Final_Proj/app.R

CONCLUSION

One limitation of our project is the completeness of the SAT scores. 6398 out of the 7703 (roughly 83%) observations have missing SAT scores, which could jeopardize the generality of our conclusion. Also, since we subsetted the data set to consider only the most recent time period (2014-2015) when average 10-year-salary is available, the observed patterns could have changed over the past few years. Additionally, since we did not cover all the features in the original data set, we might overlook some other important factors, which could include some confounding variables that intermediate the correlations we observed. Lastly, when analyzing the cost-value trade-off, since we only take continuous features into consideration, a caveat should be made regarding the categorical features such as predominant degree, which could affect the overall pattern.

Hence, future exploration could make use of larger span of the original dataset, with longer time span and more variables included. Furthermore, one could apply methods such as logistic regression to analyze the unexplored categorical features in cost-value trade-off.

REFERENCE

Shiny tutorial: <https://shiny.rstudio.com/tutorial/>

Shiny Interactive visualizations: <https://shiny.rstudio.com/gallery/movie-explorer.html>